

000

001

002

003

004

005

006

007

008

009

010

011

012

013

014

015

016

017

018

019

020

021

022

023

024

025

026

027

028

029

030

031

032

033

034

035

036

037

038

039

040

041

042

043

044

045

046

047

048

049

050

051

052

053

054

055

Secure Linear Alignment of Large Language Models

Anonymous Authors¹

Abstract

Language models increasingly appear to learn similar representations, despite differences in training objectives, architectures, and data modalities. This emerging compatibility between independently trained models introduces new opportunities for cross-model alignment to downstream objectives. Moreover, it unlocks new potential application domains, such as settings where security, privacy, or competitive constraints prohibit direct data or model sharing. In this work, we propose a privacy-preserving framework that exploits representational convergence to enable cross-silo inference between independent language models. The framework learns an affine transformation over a shared public dataset and applies homomorphic encryption to protect client queries during inference. By encrypting only the linear alignment and classification operations, the method achieves sub-second inference latency while maintaining strong security guarantees. We support this framework with an empirical investigation into representational convergence, in which we learn linear transformations between the final hidden states of independent models. We evaluate these cross-model mappings on embedding classification and out-of-distribution detection, observing minimal performance degradation across model pairs. Additionally, we show for the first time that linear alignment sometimes enables text generation across independently trained models. Our code is included in the submission and will be released upon publication.

1. Introduction

Large language models (LLMs) have become the standard paradigm for language understanding and generation, with both encoder-style and autoregressive architectures achieving strong generalization across diverse tasks (Srivastava et al., 2022). Their rapid progress is driven by scaling laws that link model size, compute, and data volume to emergent capabilities (Kaplan et al., 2020; Hoffmann et al., 2022). As LLMs continue to scale, recent evidence from the Platonic Representation Hypothesis suggests that different models are becoming more aligned in their learned representations (Huh et al., 2024), raising the possibility that independently

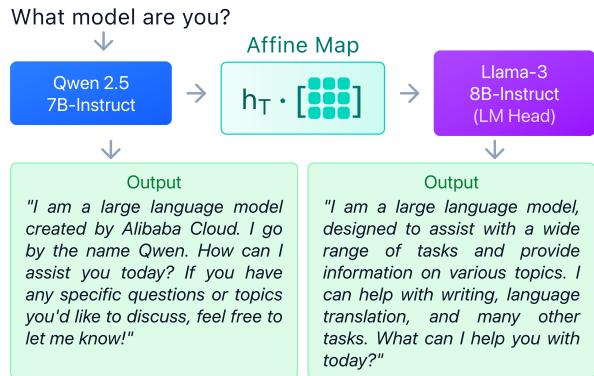


Figure 1. Text Generation via Cross-Model Linear Alignment: We learn an affine map from Qwen’s hidden states into Llama’s feature space, enabling Qwen representations to be decoded by Llama’s token head. The resulting hybrid model combines Qwen’s encoder/transformer blocks with Llama’s output head, producing coherent responses without adopting either model’s identity.

trained systems share a common latent structure. This emerging representational similarity enables new forms of model interoperability. For example, prior work on model stitching shows that independently trained models can be aligned through learned linear transformations, allowing downstream feature transfer across architectures (Bansal et al., 2021). Such compatibility supports multi-model inference pipelines, modular system design, and collaboration across heterogeneous systems (Jiang & Li, 2024; Chen et al., 2025). These capabilities are particularly valuable in settings where privacy constraints, proprietary considerations, or regulatory requirements (e.g., GDPR (Voigt & Von dem Bussche, 2017); HIPAA (Gostin & Hodge, 2000)) prohibit direct data or model sharing (Diebold, 2023; Reimsbach-Kounatze et al., 2025).

In this work we introduce **HELIX** (Homomorphically Encrypted Linear Inference across models), a privacy-preserving framework that exploits representational similarity to enable cross-silo inference between independent LLMs. The key insight is that when models learn similar representations, their feature spaces can be aligned through a simple linear map (Figure 1), which can be executed efficiently under homomorphic encryption. **HELIX** operates in two phases. During training, the client encrypts their embeddings from public data and sends them to the service provider, who computes an alignment map under encryption

and returns it. During inference, clients apply the alignment locally to their embeddings, then encrypt and send the transformed representations to the provider, who applies a linear classifier homomorphically and returns the encrypted prediction. Because the alignment and classification operations are linear, the protocol achieves 128-bit security with sub-second per-sample latency.

HELIX relies on the assumption that linear transformations can preserve task performance across models. To validate this assumption, we conduct an empirical investigation into linear alignment across diverse encoder-based and autoregressive **LLMs**. First, we verify that independently trained models exhibit nontrivial shared linear structure (Figure 2). We then measure alignment performance on supervised classification and **out-of-distribution (OOD)** detection, observing minimal degradation across models when using a fixed target linear head.

For text generation, we evaluate linear alignment across 34 model pairs using instruction-tuned generative models. Our investigation reveals two patterns: First, tokenizer compatibility strongly predicts success, with exact token match rate ($r = 0.898$) and Jaccard index ($r = 0.822$) correlating with text generation quality. Second, model scale imposes a minimum threshold, as all pairs with source models below 4B parameters produce lower quality results regardless of tokenizer alignment. We assess quality through LLM-as-a-Judge evaluation, embedding similarity to source models, and human judgment.

The remainder of this paper is organized as follows. Section 2 reviews related work on representational similarity and privacy-preserving inference. Section 3 establishes the effectiveness of linear alignment across diverse tasks. Sections 4–5 present the **HELIX** framework and experiments. In summary, our contributions are as follows:

- We propose **HELIX**, a privacy-preserving framework for cross-silo inference with independent embedding models. Leveraging the representational similarity of **LLMs**, **HELIX** encrypts only linear operations (alignment and classification) rather than full transformer models, achieving sub-second inference latency while protecting client queries.
- We provide a systematic characterization of cross-model text generation via linear alignment. Analyzing 34 **LLM** model pairs, we identify tokenizer compatibility and model size as primary determinants of success.
- We show that supervised linear alignment preserves classification and **OOD** detection performance across embedding model pairs with minimal degradation.

2. Related Work

Understanding whether neural networks converge to similar representations despite stochastic training and non-convex optimization is a central question in machine learning (Li et al., 2015; Raghu et al., 2017; Klabunde et al., 2025). A

growing body of work shows that independently trained models often discover surprisingly aligned representations, particularly in overparameterized, high-capacity regimes.

Representational Similarity. Kornblith et al. (2019) introduced **CKA**, showing that identically structured CNNs trained from different seeds learn similar intermediate features. Morcos et al. (2018) found that wider networks, and those that generalize better, exhibit stronger representational alignment. At a higher level, the Platonic Representation Hypothesis (Huh et al., 2024) proposes that large models increasingly converge toward a shared statistical understanding of the world, with similar geometric structure emerging across modalities and architectures.

Model Stitching. Building on representational convergence, model stitching asks whether a lightweight adapter can map intermediate representations from one model into another’s feature space while preserving downstream performance. Early CNN results show that shallow linear layers can stitch models trained under different settings (Lenc & Vedaldi, 2015; Bansal et al., 2021). Bansal et al. (2021) argue that stitching complements statistical similarity metrics (e.g., CKA) by testing *functional* interchangeability rather than mere geometric resemblance. Related work further shows that independently trained CNN and face-recognition models can be linearly aligned at the final layer with minimal accuracy loss (McNeely-White et al., 2020; 2022). Recent work extends stitching to transformers and **LLMs**: Chen et al. (2025) align hidden states across language models of different sizes via linear maps to transfer features. Unlike our approach, theirs focuses on computational efficiency. Jiang & Li (2024) stitch autoregressive and bidirectional transformers (GPT and BERT) for look-ahead text understanding, supporting the feasibility of cross-**LLM** alignment.

Linear Identifiability. Roeder et al. (2020) show that for a broad class of models, including supervised, contrastive, and causal language models, representations learned on the same data and architecture are *linearly identifiable*: there exists an invertible matrix W such that $Z_B \approx WZ_A$. This result provides a theoretical foundation for the use of linear alignment methods like ours. However, the identifiability theorem is guaranteed only when architectures, objectives, and data distributions match. When these conditions differ, representational equivalence may be approximate.

Homomorphic encryption. Homomorphic Encryption (HE) is widely used to protect client inputs during outsourced inference. CryptoNets (Gilad-Bachrach et al., 2016b) first demonstrated HE inference with a public model, while Gazelle (Juvekar et al., 2018) combines HE with **Secure Multi-Party Computation (MPC)** to accelerate end-to-end interactive private inference. More recent systems target secure transformer execution: Nexus (Zhang et al., 2024a) proposes non-interactive HE protocols, Powerformer (Park et al., 2024) improves throughput via HE-friendly attention, and BOLT (Pang et al., 2024) and Nimbus (Li et al., 2024)

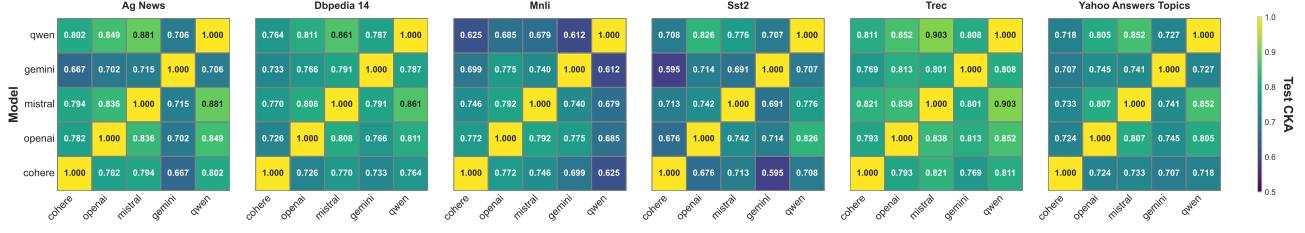


Figure 2. **Linear CKA similarity across embedding APIs.** We compute linear CKA (Kornblith et al., 2019) on vendor-provided embeddings over shared inputs from multiple datasets. CKA values range from 0.595 to 0.881, indicating substantial shared linear structure across independently trained models.

use MPC for online private inference. HETAL (Lee et al., 2023) enables encrypted fine-tuning on fixed backbones, and the Encryption-Friendly LLM architecture (Zhang et al., 2025) redesigns transformers for HE-compatible operations. All prior HE work encrypts the entire model. HELIX fundamentally differs because it extracts features locally and encrypts only the linear head.

3. Linear Alignment Characterization

In this section, we validate HELIX’s core assumption by testing whether a simple linear map can align independently trained LLMs while preserving downstream task behavior.

Preliminaries For each experiment, we consider two embedding models: the target model \mathcal{F}_A consists of a representation function $g_A : \mathcal{X} \rightarrow \mathbb{R}^{d_A}$ followed by a task head $f_A : \mathbb{R}^{d_A} \rightarrow \mathcal{Y}$. The source model \mathcal{F}_B provides a representation function $g_B : \mathcal{X} \rightarrow \mathbb{R}^{d_B}$. Throughout, $g_A(\cdot)$ and $g_B(\cdot)$ correspond to the final hidden representations produced by their respective transformer encoders.

We focus on settings where the target head (f_A) is linear. For a K -class classification task, the head takes the form $f_A(z) = zV + c$, with parameters $V \in \mathbb{R}^{d_A \times K}$ and $c \in \mathbb{R}^K$ learned on labeled training data using $g_A(x)$.

To relate the two representation spaces, we learn a linear alignment from the source to the target space. Given an input $x \in \mathcal{X}$, the aligned representation is $\hat{z}_A = g_B(x)W + b$, where $W \in \mathbb{R}^{d_B \times d_A}$ and $b \in \mathbb{R}^{d_A}$ are learned parameters. For all experiments, predictions are obtained by applying the fixed classifier $f_A(\cdot)$ to \hat{z}_A without retraining.

In our experiments, g_A and g_B are instantiated using a set of pretrained language models drawn from both vendor-provided embedding services and locally hosted autoregressive models. For embedding models, we use OpenAI’s text-embedding-3-small, Cohere’s embedding API, Google’s Gemini embedding-001 model, Qwen3-Embedding-8B, and e5-mistral-7b-instruct models. In all cases, the representation functions $g_A(\cdot)$ and $g_B(\cdot)$ are kept fixed, and only the affine alignment parameters (W, b) are learned.

Representational Similarity. Before evaluating behavioral transfer, we first verify that independently trained models exhibit shared linear structure. Figure 2 reports representative CKA similarity across model pairs; full CKA/SVCCA methodology and results are deferred to the Appendix.

Party A (Target)	Party B (Source)	Classification Acc. Baseline	Classification Acc. Lin.Map	OOD Dataset	OOD AUROC Baseline	OOD AUROC Lin.Map
SST-2 (50%)						
Gemini	OpenAI	94.5	93.1		0.826	0.801
Cohere	Gemini	94.4	91.7		0.819	0.870
OpenAI	Cohere	93.0	92.4	AGNews	0.875	0.843
Mistral	Qwen	94.5	93.7		0.859	0.826
TREC (17%)						
OpenAI	Qwen	96.0	95.6		0.738	0.916
Cohere	Gemini	97.0	91.2	AGNews	0.921	0.766
OpenAI	Cohere	96.4	94.4		0.954	0.802
Mistral	Qwen	97.0	96.6		0.921	0.940
AG News (25%)						
Gemini	OpenAI	92.6	91.6		0.908	0.891
OpenAI	Qwen	92.6	91.9	MNLI	0.953	0.913
Cohere	Gemini	91.9	91.5		0.875	0.885
Mistral	Qwen	92.6	92.4		0.953	0.940

Table 1. **Downstream performance is preserved under task supervised linear alignment:** We train a linear classifier $f_A(\cdot)$ on target embeddings $g_A(\cdot)$ (Baseline), then fit a linear map from source embeddings $g_B(\cdot)$ into the target space and evaluate using the same frozen head (Lin.Map).

3.1. Cross-Model Alignment to Downstream Tasks

We next test whether a simple linear transformation between embedding models preserves downstream performance on supervised classification and OOD detection. Since embeddings are commonly used as features for linear classifiers (Devlin et al., 2018; Tunstall et al., 2022), this provides a natural testbed for cross-model representation compatibility.

Experimental setup. For each dataset, we designate a target model (PARTY A) and a source model (PARTY B). We train the target linear classifier $f_A(\cdot)$ on the training split using target embeddings, and learn an affine map (W^*, b^*) on the same split to project source embeddings $g_B(x)$ into the target feature space. At test time, we freeze $f_A(\cdot)$ and evaluate predictions using aligned source representations:

$$\hat{y} = f_A(g_B(x)W^* + b^*).$$

OOD detection. OOD detection evaluates whether a model can separate in-distribution inputs from unseen data by probing its logits *confidence*. We use the *Energy score* (Liu et al., 2021) for logits $f_A(z) \in \mathbb{R}^K$:

$$\mathcal{E}(z) = -\log \sum_{k=1}^K \exp(f_A(z)_k).$$

165	Model 1	Model 2	Baseline Acc.		Mapped Acc.	
			M1	M2	M1→M2	M2→M1
166	Llama3-8B	Qwen2.5-7B	58	70	48	68
167	Gemma3-270M	Llama3-8B	22	58	20	49
168	Gemma3-270M	Qwen2.5-7B	22	70	21	68
169	Llama3.2-1B	Llama3-8B	42	58	36	58
170	Llama3.2-1B	Qwen2.5-7B	42	70	28	69

Table 2. **Cross-model linear alignment on MMLU:** Baseline shows native model accuracy; Mapped shows accuracy after linearly transforming Model 1’s representations to Model 2’s head ($M1 \rightarrow M2$), or vice versa. Mapping from stronger to weaker models preserves performance, while mapping from weaker to stronger models degrades substantially.

Lower Energy indicates higher confidence, while higher values are characteristic of OOD inputs. We report AUROC by thresholding $\mathcal{E}(z)$ to distinguish in- vs. out-of-distribution samples. Since Energy depends on the full logit distribution, it provides a sensitive test of whether linear alignment preserves the target model’s confidence structure.

Results. We report in-distribution classification accuracy and OOD AUROC, where OOD samples come from an alternative dataset. **Baseline** trains and evaluates a linear classifier $f_A(\cdot)$ on target embeddings $g_A(\cdot)$, while **Lin.Map** applies the same classifier to linearly aligned source embeddings. Table 1 shows that linear alignment largely preserves classification accuracy and achieves competitive OOD detection performance, with AUROC often matching or exceeding the baseline, indicating that the mapping recovers both decision boundaries and confidence structure.

3.2. Text Generation

We next evaluate whether linear alignment extends to the more demanding setting of autoregressive text generation. Unlike classification, which requires a single forward pass, generation demands coherent sequential predictions over many decoding steps where errors can accumulate, making generation a stringent test of whether linear maps preserve the structure necessary for multi-step inference.

Experimental Setup. We evaluate cross-model generation using open-source instruction-tuned models ranging from 270M to 32B parameters. For each model pair, we learn an affine transformation that maps source model B’s penultimate layer representations to target model A’s prediction head space. The transformation is learned via ordinary least squares with ridge regularization ($\lambda = 10^{-4}$) on 4,000 training examples from either MMLU or Alpaca. We extract hidden states using token-level alignment based on character offsets, matching each token from model A to the corresponding token in model B whose ending position is equal to or immediately follows A’s ending position.

Crucially, we use model A’s prediction head without fine-tuning. At each decoding step, we extract model B’s penultimate hidden state, apply the learned affine transformation to map into model A’s representation space, and pass the result to model A’s frozen LM head $f_A(\cdot)$ for next-token

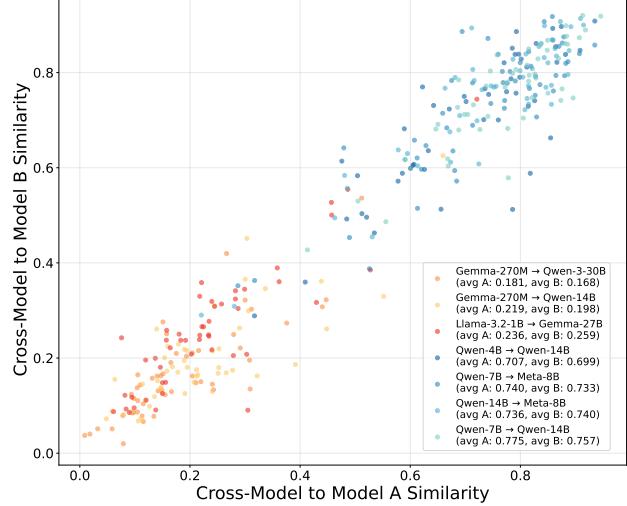


Figure 3. **Cross-Model Embedding Similarity to Native Models:** We compare Cross-Model text generation to the text produced by each base model using cosine similarity (using OpenAI’s embedding-001). Each point represents an Alpaca test prompt. High-similarity pairs (upper right) produce coherent text, while low-similarity pairs (lower left) generate incoherent outputs.

prediction. This tests whether linear mapping alone can enable zero-shot cross-model generation. We use greedy decoding for a maximum of 128 tokens.

3.2.1. MMLU: STRUCTURED OUTPUT GENERATION

Task and Evaluation. We evaluate accuracy on 100 MMLU test questions, using greedy decoding with early stopping when a valid answer letter (A, B, C, or D) is generated. For comparison, we report baseline accuracy for each model’s native performance on the same questions.

Results. Table 2 shows results for five model pairs. Mapping from stronger to weaker models preserves substantial accuracy, while mapping from weaker to stronger models degrades significantly. This asymmetric transfer demonstrates that strong representations retain semantic structure after linear transformation and can be decoded zero-shot by weaker models’ prediction heads without fine-tuning. The asymmetry confirms that source representational capacity, not the target head, is a limiting factor.

3.2.2. ALPACA: INSTRUCTION FOLLOWING

Task and Evaluation. We evaluate cross-model generation on 100 held-out Alpaca test prompts using three complementary methods: (1) embedding similarity to native model outputs using OpenAI’s text-embedding-3-small model; (2) LLM-as-a-Judge evaluation using GPT-5.1 and Claude-Sonnet-4, scoring relevance, quality, and helpfulness on a 1-10 scale; and (3) human evaluation on 200 samples (evenly split between single-model and cross-model) using a 1-5 scale based on proper English, quality, and coherence.

Results. We find that certain model pairs align substantially better than others, with two consistent trends. First,

220	Olmo - 7B → Qwen 2.5 - 14B
221	What are the benefits of social media for businesses?
222	The key benefits of using social media for businesses include increased brand awareness, improved customer engagement, enhanced advertising opportunities, better customer service, and cost-effective marketing...
223	Evaluate the following expression: $(6+2)^*(2-2)$.
224	The result of the given expression is 0. First we calculate the addition and the subtraction, which are 8 and 0, and then the 8 is multiplied by 0...
225	Describe how the given product works in 1-2 sentences: A portable charger.
226	A portable charger is a small, portable device with a high-capacity battery that recharges phones, tablets, or other electronics on the go. It has a set number of charge cycles and can recharge devices that...
227	Llama 3 - 8B → Qwen 2.5 - 14B
228	Provide advice to a person who is considering becoming an entrepreneur.
229	First, make sure you have a solid business idea. This is the foundation of your business, and it's crucial that you have a clear and well-defined idea of what you want to do. Do your research, identify your target market, and make sure you have a unique value proposition that sets your business apart from...
230	Writing a short article about a new technology that you think will revolutionize the world.
231	One new technology that has the potential to revolutionize the world is the development of a fully autonomous transportation system. This technology has been in the making for several years and is now on the verge of being fully implemented. The concept of a fully autonomous transportation system is one where vehicles are able to navigate and operate without the need for human input. This could be in the form of autonomous cars, drones...
232	
233	
234	
235	

Table 3. Chat prompts and their corresponding completions for a given model pair using cross-model linear alignment.

smaller models (≤ 1 B parameters) yield poor text generation performance when mapped to larger models, even when both models belong to the same family (Figure 3). Second, *tokenizer compatibility* is a predictor of cross-model generation success. Analyzing 23 larger pairs (≥ 4 B parameters), we find that exact token match rate (i.e. the fraction of tokens aligning at corresponding positions when tokenizing identical text) strongly correlates with generation quality ($r = 0.898, p < 0.001$; Figure 4). High-quality pairs (LLM-judge score ≥ 3.5) consistently exhibit exact match ≥ 0.67 , while failures show ≤ 0.24 . Additionally, we analyze vocabulary overlap via Jaccard index, which shows similar predictive power ($r = 0.822, p < 0.001$; Appendix C). For example, Gemma’s poor cross-family performance stems from low tokenizer compatibility (exact match < 0.23 , Jaccard < 0.07). We find that pairs with exact match > 0.7 succeed consistently, as evidenced by the clustering patterns in Figure 3. Full results are in Appendix C.

Embedding similarity. Figure 3 reveals a strong correlation between embedding similarity and generation quality. Model pairs with consistent high quality text generation cluster into either high-similarity (upper right, cosine similarity > 0.7 to both models), and model pairs with poor text generatio cluster into low-similarity (lower left, < 0.3) regions. High-similarity pairs like Qwen-7B↔Meta-8B and Qwen-14B→Meta-8B produce outputs semantically similar to both native models, while low-similarity pairs like Gemma-270M→Qwen-14B generate poor quality text.

LLM-as-a-Judge scores. Cross-model generation via linear alignment achieves LLM-judge scores of 4.0-4.7 for high-compatibility pairs (Qwen → Llama, Mistral-Nemo → Apertus), retaining 60-70% of single-model baseline. Low-compatibility pairs such as Gemma and small models (< 2 B) produce incoherent text (scores < 2.0), validating our embedding similarity analysis. Detailed scores and comparisons to single-model baselines are provided in the Appendix.

Human evaluation. Human evaluation scores corroborate



Figure 4. **Exact Token Match Rate Predicts Cross-Model Generation Quality.** Exact token match rate between two models predicts cross-model text generation quality across 23 model pairs. Quality is measured by LLM-as-a-Judge Scores ($r = 0.898, p < 0.001$).

the LLM-as-a-Judge findings. Cross-model pairs judged as low-quality by the LLM received an average human score of 1.1 (± 0.2), while high-quality cross-model pairs scored 3.0 (± 0.7). For comparison, low-quality baseline models scored 2.5 (± 0.8) and high-quality baseline models scored 4.5 (± 0.3). Human evaluators ranked outputs in the same order as the automated judge, validating that embedding similarity and LLM-as-a-Judge scores as reliable quality metrics.

Perplexity Analysis. Finally, we observe that larger model pairs (7-8B parameters) show lower perplexity degradation ($26.07\% \pm 12.26\%$) compared to smaller-to-larger mappings ($37.25\% \pm 25.84\%$), though small sample sizes ($N = 3-4$) limit statistical conclusions.

4. Privacy Framework

In this section we formalize the privacy-preserving execution model underlying HELIX, including the cryptographic assumptions, threat model, and protocol design. An extended definition of the framework is in the Appendix.

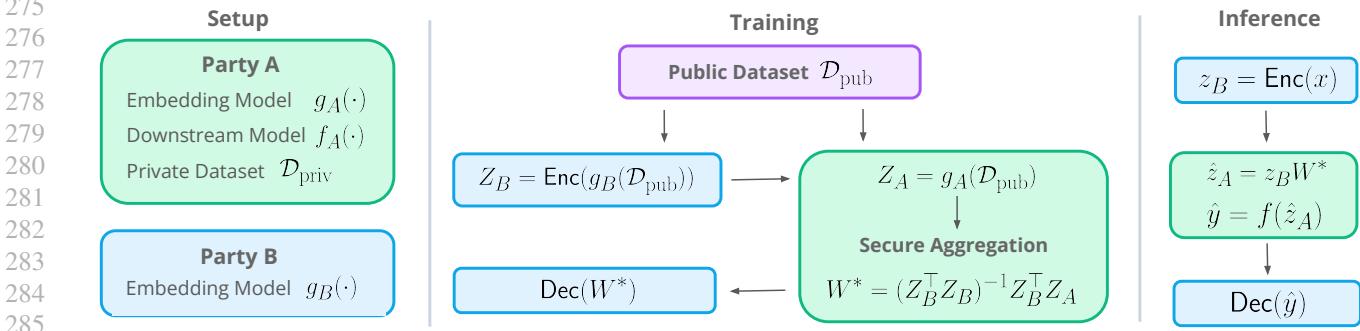


Figure 5. **Two-party privacy-preserving alignment and inference.** **Training:** PARTY B (client) encrypts embeddings $Z_B = g_B(\mathcal{D}_{\text{pub}})$ and sends $\text{Enc}(Z_B)$ to PARTY A (provider), who computes the encrypted cross-covariance $\text{Enc}(Z_A^T Z_B)$ using plaintext $Z_A = g_A(\mathcal{D}_{\text{pub}})$ and returns $\text{Enc}(Z_A^T Z_B)$ to PARTY B. PARTY B decrypts $Z_A^T Z_B$ and computes W^* locally using Eq. (1). **Inference:** PARTY B computes aligned embedding $\hat{z}_A = z_B \cdot W^* + b^*$ locally, encrypts $\text{Enc}_I(\hat{z}_A)$, and sends to PARTY A, who applies the classifier homomorphically and returns the encrypted prediction for PARTY B to decrypt.

4.1. Problem Formulation and Entities

We consider a cross-silo inference scenario involving two parties: PARTY A (service provider), who owns a classification model, and PARTY B (client), who owns an independent embedding model. PARTY A provides classification-as-a-service (e.g., an API) by returning predictions without sharing model parameters. Due to privacy, regulatory, or competitive constraints, PARTY B cannot share raw data, query embeddings, or encoder parameters with PARTY A. Our goal is to learn an affine map that aligns PARTY B’s representations to PARTY A’s feature space using encrypted communication, enabling secure inference while protecting client confidentiality.

PARTY A (Service Provider). PARTY A holds a classification head $f_A : \mathbb{R}^{d_A} \rightarrow \mathcal{Y}$ trained on embeddings from their proprietary encoder $g_A : \mathcal{X} \rightarrow \mathbb{R}^{d_A}$ using private data $\mathcal{D}_{\text{priv}}$. PARTY A provides classification-as-a-service (e.g., an API) and does not share the classifier parameters (V, c) in plaintext. Instead, PARTY A evaluates f_A server-side on encrypted inputs, enabling PARTY B to obtain predictions without revealing queries while PARTY A retains model ownership. In our setting, f_A is a linear classifier.

PARTY B (Client). PARTY B holds a proprietary embedding model $g_B : \mathcal{X} \rightarrow \mathbb{R}^{d_B}$. PARTY B seeks to leverage PARTY A’s classifier f_A for downstream predictions without revealing query data, embeddings $g_B(x)$, or encoder parameters to PARTY A.

Encrypted Computation. Our protocol uses homomorphic encryption (CKKS (Cheon et al., 2017)) to enable computation on encrypted data without revealing inputs. We restrict secure computation to linear operations, as approximate HE schemes evaluate these efficiently, following prior work that targets linear components for encrypted training and inference (Gilad-Bachrach et al., 2016a; Mohassel & Zhang, 2017; Juvekar et al., 2018; Lee et al., 2023). Our implementation returns full logits, however, prior work has shown that encrypted argmax operations (Jovanovic et al., 2022; Zhang et al., 2024a) can limit information leakage

and defend against model extraction and membership inference attacks (Tramèr et al., 2016; Carlini et al., 2021). These techniques could be straightforwardly integrated with our linear alignment protocol for enhanced model provider privacy in production deployments.

4.2. Threat Model and Guarantees

We adopt a *semi-honest* (honest-but-curious) threat model (Goldreich, 2004) where both parties follow the protocol but may analyze communication to infer private information; we do not consider malicious adversaries. We prioritize client-side privacy with the following security guarantees:

- **Client input privacy (PARTY B).** During inference, PARTY A observes only CKKS ciphertexts (encrypted aligned embeddings and encrypted outputs). Without PARTY B’s secret key, semantic security implies these ciphertexts reveal no information about queries x , embeddings $g_B(x)$, or encoder parameters beyond what is implied by the decrypted final predictions.
- **Provider classifier privacy (PARTY A).** PARTY A’s classifier parameters (V, c) are never revealed in plaintext to PARTY B and are only used within encrypted linear inference. The base protocol returns encrypted prediction outputs to PARTY B. We treat adaptive attacks such as model extraction and membership inference as *out of scope* for this work.

Alignment map leakage. PARTY B retains the learned alignment map W^* in plaintext after decryption and uses it locally during inference. While W^* reveals structural information about PARTY A’s embedding space (dimension d_A), it does not expose the classifier parameters (V, c) or training data $\mathcal{D}_{\text{priv}}$. However, W^* may enable property inference about PARTY A’s model architecture. We evaluate membership inference attacks on W^* in Appendix J, demonstrating that W^* does not leak individual sample membership under standard geometric feature extraction.

Limitations. As in most HE deployments, structural information (embedding dimensions, communication volume)

may be revealed. Our semi-honest threat model assumes PARTY B makes honest inference queries rather than attempting model extraction through adaptive attacks.

4.3. Secure Linear Alignment

Given public data \mathcal{D}_{pub} with embeddings $Z_A = g_A(\mathcal{D}_{\text{pub}})$ and $Z_B = g_B(\mathcal{D}_{\text{pub}})$, we compute the affine alignment from the sufficient statistics $Z_B^\top Z_B$ and $Z_A^\top Z_A$:

$$W^* = (Z_B^\top Z_B + \lambda I)^{-1} Z_B^\top Z_A. \quad (1)$$

To avoid revealing raw embeddings, we stream aggregates over mini-batches; the cross-covariance $Z_B^\top Z_A$ is computed under homomorphic encryption via secure aggregation.

Public data assumption. We assume a shared, non-sensitive public dataset \mathcal{D}_{pub} for fitting W^* . When learned exclusively from \mathcal{D}_{pub} , the alignment reflects only public distributional structure and PARTY A’s private data remains confidential. Optionally, PARTY A may augment \mathcal{D}_{pub} with a small number of plaintext in-distribution samples from $\mathcal{D}_{\text{priv}}$ (64-128 samples) to improve alignment accuracy; this trades data confidentiality for improved performance, revealing limited task-specific information to PARTY B. While the shared samples are directly exposed, a secondary concern is whether W^* itself leaks information about which specific samples were included—for instance, if PARTY A curates the shared set to exclude sensitive examples, can an adversary infer their presence or absence from W^* ? We evaluate this through membership inference analysis in Appendix J, demonstrating that W^* leakage is provably bounded: membership inference advantage is $O(\sqrt{d}/N) \approx 0.016$ for our configuration, yielding negligible privacy risk.

4.4. Two-Party Secure Training Protocol

To compute Eq. (1) while preserving client confidentiality:

1. **Encrypt and send.** PARTY B (client) generates keys (pk, sk) , encrypts Z_B to obtain $\text{Enc}_{pk}(Z_B)$, and sends it to PARTY A.
2. **Secure aggregation.** PARTY A (provider) computes $\text{Enc}(Z_A^\top Z_B)$ via homomorphic matrix multiplication using plaintext Z_A and encrypted Z_B , then returns $\text{Enc}(Z_A^\top Z_B)$ to PARTY B.
3. **Decrypt and solve.** PARTY B decrypts to obtain $Z_A^\top Z_B$, then computes $W^* = (Z_B^\top Z_B + \lambda I)^{-1} Z_A^\top Z_B$ in plaintext using their local Z_B . Optionally, PARTY B computes bias $b^* \leftarrow \frac{1}{N} \mathbf{1}^\top (Z_A - Z_B W^*)$ after obtaining Z_A via an additional encrypted transmission or using public data statistics.

4.5. Privacy-Preserving Inference

Inference uses the plaintext map W^* held by PARTY B (client). To preserve query confidentiality, PARTY B generates a fresh inference keypair (pk_I, sk_I) and shares pk_I with PARTY A. Assuming a linear classifier $f_A(z) = zV + c$, the protocol is as follows:

1. **Compute and encrypt aligned representation.**

PARTY B computes $z_B = g_B(x)$, applies the alignment locally $\hat{z}_A = z_B \cdot W^* + b^*$, and sends $\text{Enc}_{pk_I}(\hat{z}_A)$ to PARTY A.

2. **Homomorphic classification.** PARTY A computes $\text{Enc}(\hat{y}) = \text{Enc}(\hat{z}_A) \cdot V + c$ homomorphically.
3. **Decrypt output.** PARTY B decrypts $y = \text{Dec}_{sk_I}(\text{Enc}(\hat{y}))$.

5. Experiments

We apply the [HELIX](#) framework to large-scale embedding models to evaluate its effectiveness on downstream tasks. Embeddings are commonly used directly as features to train a linear classifier on labeled data ([Devlin et al., 2018](#); [Wang et al., 2022](#)). Our experiments utilize five embedding models: EMBEDDING-001 (Google), TEXT-EMBEDDING-3-SMALL (OpenAI), E5-MISTRAL-7B-INSTRUCT, QWEN3-EMBEDDING-8B, and EMBED-ENGLISH-V3.0 (Cohere).

5.1. Embedding Classification

We investigate the cross-silo knowledge transfer capabilities of [HELIX](#) between two private embedding models. In contrast to Section 3, we concentrate on training a linear map on a public dataset, rather than the in-distribution dataset.

Experimental Setup. We choose an embedding model to represent PARTY A and train a linear classifier $f(\cdot)$ on the full training set using embeddings $Z_A = g_A(X)$ and labels Y . We evaluate across six datasets: TREC, MNLI, DBpedia, and AG News. We use embedding model $g_B(\cdot)$ to serve as PARTY B.

We then train the linear map W^* under two settings: (1) using only a public, *independent* dataset to estimate W^* , and (2) augmenting the public dataset with a small number of in-distribution samples (64, 128) in the dataset under test. For the public dataset, we use Wikipedia and IMDB. Our baseline are a linear classifier trained on the full training set ($f(\cdot)$), as well as a classifier trained on the few-shot in-distribution samples (in practice the client would have access to the few-shot samples to build their own classifier). Setting (2) serves as an upper-bound analysis demonstrating the performance ceiling when W^* uses limited in-distribution data (64-128 samples). However, this compromises data security by requiring PARTY A to share potentially sensitive samples with PARTY B during training, violating the zero-shot privacy guarantees of Setting (1). We argue that this performance gap quantifies the privacy-utility trade-off: how much accuracy is sacrificed to preserve full data confidentiality versus accepting limited data exposure for improved alignment.

Results. Table 4 shows that cross-model alignment achieves strong performance across model pairs and datasets. **Public + ID** (64-128 samples) consistently matches or exceeds baselines trained on the same few-shot data, with particularly strong results on TREC and MNLI. Critically, **Public Only** surpasses 64-shot baselines in many configurations.

Party A (Target)	Party B (Source)	Baseline Full Data (%)	Baseline 64	Baseline 128	Public Only	Public + ID 64	Public + ID 128	Baseline Full Data (%)	Baseline 64	Baseline 128	Public Only	Public + ID 64	Public + ID 128
TREC (17%)													
MNL (33%)													
Gemini	OpenAI	95.4	56.6	77.6	57.6	78.0	81.0	65.0	33.5	35.1	44.5	48.5	48.9
OpenAI	Qwen	96.4	59.0	<u>82.8</u>	58.4	77.6	<u>83.6</u>	62.4	57.7	61.9	64.0	<u>64.9</u>	65.2
Mistral	Cohere	96.6	52.0	73.2	55.8	<u>74.4</u>	78.4	77.6	34.4	35.8	39.5	<u>41.9</u>	42.8
OpenAI	Cohere	96.4	52.2	73.4	68.4	80.2	<u>80.6</u>	62.4	34.4	35.8	46.5	<u>47.1</u>	47.9
Mistral	Qwen	96.6	59.4	82.2	65.6	83.8	<u>87.4</u>	77.6	57.6	61.8	58.8	<u>71.3</u>	72.9
Qwen	OpenAI	97.0	57.0	77.8	75.4	82.8	<u>87.6</u>	87.9	33.4	35.2	42.0	42.9	43.9
DBpedia (7%)													
Gemini	OpenAI	99.1	65.0	77.8	53.5	<u>89.3</u>	94.1	92.1	79.6	87.0	85.9	<u>88.4</u>	88.7
OpenAI	Qwen	99.1	66.0	83.4	66.7	<u>91.0</u>	95.8	92.7	80.0	87.1	79.7	<u>88.6</u>	89.2
Mistral	Cohere	99.2	56.5	<u>78.9</u>	44.5	70.0	<u>86.7</u>	93.4	81.8	85.0	64.3	80.0	<u>83.6</u>
Gemini	Mistral	99.1	52.0	79.1	65.7	<u>93.4</u>	96.6	92.1	75.7	87.8	86.7	<u>88.1</u>	88.7
Mistral	Qwen	99.2	72.9	93.2	79.6	<u>93.6</u>	96.5	93.4	84.4	86.5	71.4	<u>87.9</u>	88.2
Qwen	OpenAI	99.1	74.1	<u>88.9</u>	66.4	81.3	89.9	92.9	82.0	86.1	84.6	<u>86.2</u>	87.2
AG News (25%)													

Table 4. **Cross-model alignment classification.** A classifier $f(\cdot)$ is trained on private data (column *Baseline Full Data (%)*). We evaluate **HELIX**s by mapping representations from a public dataset, and from the public dataset augmented with few-shot in-distribution samples (64, 128), prior to classification by $f(\cdot)$. Additional baselines correspond to independently trained classifiers using *only* the few-shot samples. Best results are shown in **bold**, second-best underlined. Each experiment is repeated three times with different random seeds. The standard deviation across runs is ± 1.5 percentage points, indicating high reproducibility. Full results are in the Appendix.

Method	Security Scope	SST-2	MRPC	RTE	Latency (s)
Baseline	None	92.3	90.3	69.7	< 1
BOLT	Full (HE+MPC)	92.8	90.0	69.3	> 60
Nimbus	Full (HE+MPC)	92.6	89.8	66.8	> 20
MPCFormer	Full (MPC)	–	88.7	64.9	18
Enc.-Friendly	Full (HE)	81.9	81.5	59.3	26.5
PowerFormer	Full (HE)	92.0	87.8	69.8	> 20
Nexus	Full (HE)	92.1	–	69.9	37.3
HELIX* (Base)	Linear (HE)	92.3	77.8	59.6	< 1
HELIX (FT)	Linear (HE)	93.0	82.0	55.6	< 1

Table 5. Efficiency–utility comparison across *distinct secure inference settings* on GLUE tasks. **Security Scope** denotes which model components are evaluated under cryptographic protection. Prior methods secure full transformer inference, whereas **HELIX** encrypts only linear alignment and classification.

5.2. Computational Efficiency Analysis

We contextualize **HELIX** alongside prior private inference systems on GLUE tasks in a two-party setting, while emphasizing that these approaches target **different secure-inference scopes**. Prior systems protect end-to-end transformer inference under cryptographic protocols, whereas **HELIX** assumes the client can compute embeddings locally and requires secure computation only for the final linear classification head. We report task accuracy and end-to-end latency to characterize the tradeoff within each scope.

In addition, the evaluated model families differ across settings: prior work reports secure inference results for a single BERT model, while **HELIX** is assessed using randomly paired embedding models from Gemini, OpenAI, and Cohere, as well as a fine-tuned Llama-2-8b mapped to base Llama-2-8b. These experiments reflect **HELIX**’s design goal of enabling cross-silo transfer across independent models with minimal encrypted computation.

Experimental Setup. We evaluate **HELIX** on SST-2, MRPC, and RTE. Reported latency includes plaintext embedding inference, CKKS encryption, encrypted evaluation of the linear classification head, and decryption. Baselines perform cryptographically protected end-to-end inference

on BERT, as reported in their original works. Our results are averaged over three model pairs per task, with a low variance across runs. We implement **HELIX** using TenSEAL CKKS with `poly_modulus_degree=8192`, coefficient modulus [60, 40, 40, 60], and scale 2^{40} . These parameters achieve 128-bit security according to the Homomorphic Encryption Security Standard ([HomomorphicEncryption.org, 2018](#)). Our circuit has minimal multiplicative depth (depth-1: one ciphertext–plaintext multiplication), well within CKKS’s noise budget and requiring no bootstrapping or modulus switching beyond standard rescaling. All measurements are run on CPU and reported as end-to-end latency.

Results. Across tasks, **HELIX** achieves strong accuracy with sub-second latency by encrypting only the final linear head. In contrast, prior systems incur substantially higher cryptographic overhead because they secure end-to-end transformer inference (often requiring online interaction), leading to multi-second to minute-scale latency. In addition to low-latency inference, **HELIX** incurs less than 1MB of communication per sample, since it transmits only a single embedding vector (typically 1536–3072 dimensions) since CKKS uses SIMD packing to encrypt multiple values into a single ciphertext. With `poly_modulus_degree=8192`, each ciphertext can encode up to 4096 values.

6. Conclusion

We introduce a framework for linearly and securely transferring knowledge between **LLMs**. Several areas remain for future work. First, extending to multi-modal models and developing tokenizer-agnostic alignment methods could broaden applicability beyond text-only models with compatible vocabularies. Second, while our protocol protects client queries, better securing the provider’s model parameters remains an open challenge. Finally, exploring efficient non-linear alignment methods could improve cross-model generation quality while maintaining computational efficiency.

440 Impact Statement

441 This work advances the practical understanding of representation alignment and secure computation for machine
 442 learning models. The proposed techniques enable cross-
 443 model knowledge transfer and privacy-preserving training
 444 and inference in settings where data or model sharing is
 445 restricted. We do not anticipate new societal risks beyond
 446 those commonly associated with the deployment of large
 447 language models; however, as with any enabling technology,
 448 these methods could be misused depending on the
 449 application context. We therefore encourage responsible
 450 deployment consistent with established privacy, security,
 451 and ethical guidelines.

452 References

453 Ainsworth, S. K., Hayase, J., and Srinivasa, S. Git re-basin:
 454 Merging models modulo permutation symmetries. *arXiv*
 455 preprint arXiv:2209.04836, 2022.

456 Bansal, Y., Nakkiran, P., and Barak, B. Revisiting model
 457 stitching to compare neural representations. In *Advances*
 458 in *Neural Information Processing Systems (NeurIPS 34)*,
 459 pp. 225–236, 2021.

460 Brakerski, Z., Gentry, C., and Vaikuntanathan, V. Lev-
 461 eled fully homomorphic encryption without bootstrap-
 462 ping. *ACM Trans. Comput. Theory*, 6(3), 2014.

463 Carlini, N., Tramèr, F., Wallace, E., Jagielski, M., Herbert-
 464 Voss, A., Lee, K., Roberts, A., Brown, T., Song, D., Er-
 465 lingsson, Ú., Oprea, A., and Raffel, C. Extracting training
 466 data from large language models. In *30th USENIX Security*
 467 Symposium (USENIX Security 21), pp. 2633–2650.
 468 USENIX Association, August 2021. ISBN 978-1-939133-
 469 24-3. URL <https://www.usenix.org/conference/usenixsecurity21/presentation/carlini-extracting>.

470 Carlini, N., Paleka, D., Dj Dvijotham, K., Steinke, T.,
 471 Hayase, J., Cooper, A. F., Lee, K., Jagielski, M., Nasr,
 472 M., Conmy, A., Yona, I., Wallace, E., Rolnick, D., and
 473 Tramèr, F. Stealing part of a production language model,
 474 2024. URL <https://arxiv.org/abs/2403.06634>.

475 Chaudhuri, K., Monteleoni, C., and Sarwate, A. D. Differ-
 476 entially private empirical risk minimization. *Journal of*
 477 *Machine Learning Research*, 12(3), 2011.

478 Chen, A., Merullo, J., Stolfo, A., and Pavlick, E. Transfer-
 479 ing features across language models with model stitching.
 480 *arXiv* preprint arXiv:2506.06609, 2025.

481 Cheon, J. H., Kim, A., Kim, M., and Song, Y. Homomor-
 482 phic encryption for arithmetic of approximate numbers.

483 In *International conference on the theory and applica-*
 484 *tion of cryptology and information security*, pp. 409–437.
 485 Springer, 2017.

486 Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert:
 487 Pre-training of deep bidirectional transformers for lan-
 488 guage understanding. *arXiv:1810.04805*, 2018.

489 Diebold, G. Overcoming barriers to data sharing in the
 490 united states. Technical report, Center for Data Innova-
 491 tion, 2023. URL <https://www2.datainnovation.org/2023-data-sharing-barriers.pdf>.
 492 Report on legal, social, technical and economic obstacles
 493 to data sharing.

494 Entezari, R., Sedghi, H., Saukh, O., and Neyshabur, B. The
 495 role of permutation invariance in linear mode connectivity
 496 of neural networks. *arXiv preprint arXiv:2110.06296*,
 497 2021.

498 et al., K. B. Practical secure aggregation for privacy-
 499 preserving machine learning. In *ACM CCS*, 2017.

500 et al., P. R. A generic framework for privacy preserving
 501 deep learning. *arXiv preprint arXiv:1811.04017*, 2018.

502 Furlanello, T., Lipton, Z. C., Tschanne, M., Itti, L.,
 503 and Anandkumar, A. Born-again neural networks.
 504 *arXiv:1805.04770*, 2018.

505 Gentry, C. *A Fully Homomorphic Encryption Scheme*. Stan-
 506 ford University, 2009.

507 Gilad-Bachrach, R., Dowlin, N., Laine, K., Lauter, K.,
 508 Naehrig, M., and Wernsing, J. Cryptonets: Applying
 509 neural networks to encrypted data with high throughput
 510 and accuracy. In *Proceedings of the 33rd International*
 511 *Conference on Machine Learning (ICML)*, pp. 201–210,
 512 2016a.

513 Gilad-Bachrach, R., Dowlin, N., Laine, K., Lauter, K.,
 514 Naehrig, M., and Wernsing, J. Cryptonets: Applying
 515 neural networks to encrypted data with high throughput
 516 and accuracy. In *Proceedings of the 33rd International*
 517 *Conference on Machine Learning (ICML)*, 2016b.

518 Goldreich, O. *Foundations of Cryptography: Volume 2*.
 519 Cambridge University Press, 2004.

520 Gostin, L. O. and Hodge, J. G. Personal privacy and
 521 the health information portability and accountability act
 522 (hipaa): A comment on the supreme court’s decision in
 523 Ferguson v. City of Charleston. *Journal of Law, Medicine*
 524 & Ethics, 28(2):210–213, 2000.

525 Hao, M., Li, H., Chen, H., Xing, P., Xu, G., and Zhang,
 526 T. Iron: Private inference on transformers. *Advances in*
 527 *neural information processing systems*, 35:15718–15731,
 528 2022.

- 495 Hendrycks, D. and Gimpel, K. A baseline for detecting
 496 misclassified and out-of-distribution examples in neural
 497 networks. *arXiv preprint arXiv:1610.02136*, 2016.
- 498 Hinton, G., Vinyals, O., and Dean, J. Distilling the knowl-
 499 edge in a neural network. In *NIPS Deep Learning Work-
 500 shop*, 2015.
- 502 Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E.,
 503 Cai, T., Rutherford, E., Casas, D. d. L., Hendricks, L. A.,
 504 Welbl, J., Clark, A., et al. Training compute-optimal
 505 large language models. *arXiv preprint arXiv:2203.15556*,
 506 2022.
- 507 HomomorphicEncryption.org. Homomorphic encryption
 508 security standard. Technical report, HomomorphicEn-
 509 cryption.org, November 2018. URL <https://homomorphicencryption.org/standard/>. Tech-
 510 nical Report.
- 511 Huh, M., Cheung, B., Wang, T., and Isola, P. Position:
 512 The platonic representation hypothesis. In *Proceed-
 513 ings of the 41st International Conference on Machine
 514 Learning*, volume 235 of *Proceedings of Machine Learn-
 515 ing Research*, pp. 20617–20642. PMLR, 2024. URL
<https://arxiv.org/abs/2405.07987>.
- 516 Jiang, F. and Li, J. Stitching gpt to bert for look-ahead
 517 language understanding. In *Proceedings of the 62nd
 518 Annual Meeting of the Association for Computational
 519 Linguistics (ACL)*, 2024.
- 520 Jovanovic, N., Fischer, M., Steffen, S., and Vechev, M.
 521 Private and reliable neural network inference. In *Proceed-
 522 ings of the 2022 ACM SIGSAC Conference on Computer
 523 and Communications Security*, pp. 1663–1677, 2022.
- 524 Juvekar, C., Vaikuntanathan, V., and Chandrakasan, A.
 525 Gazelle: A low latency framework for secure neural net-
 526 work inference. In *USENIX Security*, 2018.
- 527 Kairouz, P., McMahan, H. B., Avent, B., Bellet, A., Bennis,
 528 M., Bhagoji, A. N., and et al. Advances and open prob-
 529 lems in federated learning. *Foundations and Trends in
 530 Machine Learning*, 14(1–2):1–210, 2021.
- 531 Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B.,
 532 Chowdhery, A., Chen, R., Elhage, N., Ebrahimi, R., Mus-
 533 abi, R., Khan, G., et al. Scaling Laws for Neural Lan-
 534 guage Models. *arXiv preprint arXiv:2001.08361*, 2020.
- 535 Karimireddy, S. P., Kale, S., Mohri, M., Stich, S., and
 536 Suresh, A. T. Scaffold: Stochastic controlled averag-
 537 ing for federated learning. In *Proceedings of the 37th
 538 International Conference on Machine Learning (ICML)*,
 539 pp. 5132–5143, 2020.
- 540 Kim, Y. and Rush, A. M. Sequence-level knowledge distil-
 541 lation. In *EMNLP*, pp. 1317–1327, 2016.
- 542 Klabunde, M., Schumacher, T., Strohmaier, M., and Lem-
 543 merich, F. Similarity of neural network models: A survey
 544 of functional and representational measures. *ACM Com-
 545 puting Surveys*, 57(9):1–52, 2025.
- 546 Knott, B., Venkataraman, S., Hannun, A., Sengupta, S.,
 547 Ibrahim, M., and van der Maaten, L. Crypten: Secure
 548 multi-party computation meets machine learning. In
 549 *NeurIPS*, 2021. URL <https://papers.neurips.cc/paper/2021/file/2754518221cfbc8d25c13a06a4cb8421-Paper.pdf>.
- 550 Kornblith, S., Norouzi, M., Lee, H., and Hinton, G. Sim-
 551 ilarity of neural network representations revisited. In
 552 *International conference on machine learning*, pp. 3519–
 553 3529. PMIR, 2019.
- 554 Lee, S., Lee, G., Kim, J. W., Shin, J., and Lee, M.-K. Hetal:
 555 Efficient privacy-preserving transfer learning with ho-
 556 momorphic encryption. In *International conference on
 557 machine learning*, pp. 19010–19035. PMLR, 2023.
- 558 Lenc, K. and Vedaldi, A. Understanding image representa-
 559 tions by measuring their equivariance and equivalence. In
 560 *Proceedings of the IEEE Conference on Computer Vision
 561 and Pattern Recognition (CVPR)*, 2015.
- 562 Li, D., Wang, H., Shao, R., Guo, H., Xing, E., and Zhang, H.
 563 MPCFORMER: FAST, PERFORMANT AND PRIVATE
 564 TRANSFORMER INFERENCE WITH MPC. In *The
 565 Eleventh International Conference on Learning Repre-
 566 sentations*, 2023. URL <https://openreview.net/forum?id=CWmvjOEhgH->.
- 567 Li, T., Sahu, A. K., Talwalkar, A., and Smith, V. Federated
 568 optimization in heterogeneous networks. In *Proceedings
 569 of Machine Learning and Systems (MLSys)*, 2020.
- 570 Li, Y., Yosinski, J., Clune, J., Lipson, H., and Hopcroft,
 571 J. Convergent learning: Do different neural net-
 572 works learn the same representations? *arXiv preprint
 573 arXiv:1511.07543*, 2015.
- 574 Li, Z., Yang, K., Tan, J., Lu, W.-j., Wu, H., Wang, X., Yu,
 575 Y., Zhao, D., Zheng, Y., Guo, M., and Leng, J. Nimbus:
 576 Secure and efficient two-party inference for transformers.
 577 In *Advances in Neural Information Processing Systems
 578 (NeurIPS 2024)*, 2024. URL <https://arxiv.org/abs/2411.15707>. arXiv:2411.15707.
- 579 Liu, W. et al. Energy-based out-of-distribution detection. In
 580 *ICLR*, 2021.
- 581 Liu, Y., Kang, Y., Xing, C., Chen, T., and Yang, Q. A secure
 582 federated transfer learning framework. *IEEE Intelligent
 583 Systems*, 2024.

- 550 Systems, 35(4):70–82, 2020. URL <https://arxiv.org/pdf/1812.03337.pdf>.
- 551
- 552 Liu, Y., Chen, T., and Yang, Q. Hierarchical federated
553 learning: Algorithms, applications and challenges. *arXiv*
554 preprint arXiv:2109.11846, 2022.
- 555
- 556 Luo, J., Zhang, Y., Zhang, Z., Zhang, J., Mu, X., Wang,
557 H., Yu, Y., and Xu, Z. Secformer: Fast and accurate
558 privacy-preserving inference for transformer models via
559 smpc. *arXiv preprint arXiv:2401.00793*, 2024. URL
560 <https://arxiv.org/abs/2401.00793>.
- 561
- 562 McMahan, H. B., Moore, E., Ramage, D., Hampson, S.,
563 and y Arcas, B. A. Communication-efficient learning of
564 deep networks from decentralized data. In *Proceedings*
565 of *Machine Learning Research (AISTATS)*, volume 54,
566 pp. 1273–1282, 2017. URL <https://proceedings.mlr.press/v54/mcmahan17a/mcmahan17a.pdf>.
- 567
- 568 McNeely-White, D., Beveridge, J. R., and Draper, B. A. In-
569 ception and resnet features are (almost) equivalent. *Cog-
570 nitive Systems Research*, 59:312–318, 2020.
- 571
- 572 McNeely-White, D., Sattelberg, B., Blanchard, N., and Bev-
573 eridge, R. Canonical face embeddings. *IEEE Transac-
574 tions on Biometrics, Behavior, and Identity Science*, 4(2):
575 197–209, 2022.
- 576
- 577 Mohassel, P. and Zhang, Y. Secureml: A system for scalable
578 privacy-preserving machine learning. In *IEEE Sym-
579 posium on Security and Privacy (SP)*, pp. 19–38, 2017. URL
580 <https://www.ieee-security.org/TC/SP2017/papers/466.pdf>.
- 581
- 582 Mora, A., Tenison, I., Bellavista, P., and Rish, I. Knowledge
583 distillation for federated learning: a practical guide. *arXiv*
584 preprint arXiv:2211.04742, 2022.
- 585
- 586 Morcos, A. S., Raghu, M., and Bengio, S. Insights on repre-
587 sentational similarity in neural networks with canonical
588 correlation. In *Advances in Neural Information Process-
589 ing Systems*, volume 31, 2018.
- 590
- 591 Pan, S. J. and Yang, Q. A survey on transfer learning. *IEEE*
592 *Transactions on Knowledge and Data Engineering*, 2010.
- 593
- 594 Pang, Q., Zhu, J., Möllering, H., Zheng, W., and Schnei-
595 der, T. Bolt: Privacy-preserving, accurate and efficient
596 inference for transformers. In *2024 IEEE Symposium on*
597 *Security and Privacy (SP)*, pp. 4753–4771. IEEE, 2024.
- 598
- 599 Park, D., Lee, E., and Lee, J.-W. Powerformer: Effi-
600 cient privacy-preserving transformer with batch rectifier-
601 power max function and optimized homomorphic atten-
602 tion. Technical report, Cryptology ePrint Archive, 2024.
- 603
- 604 Raghu, M., Gilmer, J., Yosinski, J., and Sohl-Dickstein,
605 J. Svcca: Singular vector canonical correlation analysis
606 for deep understanding and improvement. *stat*, 1050:19,
607 2017.
- 608
- 609 Reimsbach-Kounatze, C., Ishikawa, S., et al. Sharing
610 trustworthy ai models with privacy-enhancing technolo-
611 gies. Technical Report No. 38, OECD, 2025. URL
612 https://www.oecd.org/content/dam/oecd/en/publications/reports/2025/06/sharing-trustworthy-ai-models-with-privacy-enhancing-technologies_5df6fd05/a266160b-en.pdf.
- 613
- 614 Roeder, G., Wu, Y., Duvenaud, D., and Grosse, R. On linear
615 identifiability of learned representations. *arXiv preprint*
616 arXiv:2007.00810, 2020.
- 617
- 618 Romero, A., Ballas, N., Kahou, S. E., Chassang, A., Gatta,
619 C., and Bengio, Y. Fitnets: Hints for thin deep nets. In
620 *ICLR*, 2015.
- 621
- 622 Smith, V., Chiang, C.-K., Sanjabi, M., and Talwalkar,
623 A. Federated multi-task learning. *arXiv preprint*
624 arXiv:1705.10467, 2017.
- 625
- 626 Srivastava, A., Rastogi, A., Rao, A., Shoeb, A. A. M., Abid,
627 A., Fisch, A., Brown, A. R., Santoro, A., Gupta, A.,
628 Garriga-Alonso, A., et al. Beyond the imitation game:
629 Quantifying and extrapolating the capabilities of language
630 models. *arXiv preprint arXiv:2206.04615*, 2022. URL
631 <https://arxiv.org/abs/2206.04615>.
- 632
- 633 Tang, Y. and Yang, Y. Pooling and attention: What are
634 effective designs for llm-based embedding models? *arXiv*
635 preprint arXiv:2409.02727, 2024.
- 636
- 637 Tian, Y., Krishnan, D., and Isola, P. Contrastive representa-
638 tion distillation. In *ICLR*, 2020.
- 639
- 640 Tramèr, F., Zhang, F., Juels, A., Reiter, M. K., and Risten-
641 part, T. Stealing machine learning models via prediction
642 apis. In *Proceedings of the 25th USENIX Conference on*
643 *Security Symposium, SEC’16*, pp. 601–618, USA, 2016.
644 USENIX Association. ISBN 9781931971324.
- 645
- 646 Tunstall, L., Reimers, N., Jo, U. E. S., Bates, L., Korat, D.,
647 Wasserblat, M., and Pereg, O. Efficient few-shot learning
648 without prompts. *arXiv preprint arXiv:2209.11055*, 2022.
- 649
- 650 Vepakomma, P., Gupta, O., Swedish, T., and Raskar, R. Split
651 learning for health: Distributed deep learning without
652 sharing raw patient data. In *ICLR 2019 Workshop on AI*
653 for *Social Good*, 2019. URL https://aiforsocialgood.github.io/iclr2019/accepted/track1/pdfs/31_aisg_iclr2019.pdf.

605 Voigt, P. and Von dem Bussche, A. The european general
606 data protection regulation (gdpr). *A Practical Guide*.
607 Springer International Publishing, pp. 10–15, 2017.

608
609 Wang, J., Chen, Q., Sun, H., Shi, Z., and Yang, E. a. Tack-
610 ling the objective inconsistency problem in heterogeneous
611 federated optimization. *Advances in Neural Information
612 Processing Systems*, 33:7611–7623, 2020.

613 Wang, L., Yang, N., Huang, X., Jiao, B., Yang, L.,
614 Jiang, D., Majumder, R., and Wei, F. Text embeddings
615 by weakly-supervised contrastive pre-training. *arXiv
616 preprint arXiv:2212.03533*, 2022.
617

618 Wu, F., Cui, L., Yao, S., and Yu, S. Inference attacks in
619 machine learning as a service: a taxonomy, review, and
620 promising directions. *arXiv e-prints*, pp. arXiv–2406,
621 2024.

622 Zhang, J., Yang, X., He, L., Chen, K., Lu, W.-j., Wang,
623 Y., Hou, X., Liu, J., Ren, K., and Yang, X. Secure trans-
624 former inference made non-interactive. *Cryptology ePrint
625 Archive*, 2024a.

626 Zhang, L., Li, M. Y., and Griffiths, T. L. What should embed-
627 dings embed? autoregressive models represent latent gen-
628 erating distributions. *arXiv preprint arXiv:2406.03707*,
629 2024b.

630 Zhang, Y., Wang, H., et al. Encryption-friendly large lan-
631 guage model architectures. In *International Conference
632 on Learning Representations (ICLR)*, 2025.

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647

648

649

650

651

652

653

654

655

656

657

658

659

660 Appendix

661 A. LLM Similarity

662 To assess the extent to which independently trained embedding models learn compatible linear structure, we analyze
 663 representational similarity using two complementary metrics: linear Centered Kernel Alignment (CKA) and Singular Vector
 664 Canonical Correlation Analysis (SVCCA). Both measures operate directly on representation matrices computed over a
 665 shared set of inputs and quantify the degree of shared linear geometry across models.
 666

667 A.1. CKA Heatmaps

668 Our primary similarity analysis uses linear **CKA** (Kornblith et al., 2019), which measures similarity between two representa-
 669 tion matrices $Z_A \in \mathbb{R}^{n \times d_1}$ and $Z_B \in \mathbb{R}^{n \times d_2}$ computed over the same set of n inputs. Linear CKA is invariant to isotropic
 670 scaling and orthogonal transformations, making it a stable and widely used metric for comparing internal representations
 671 across models with different dimensionalities or parameterizations.
 672

Formally, linear CKA is defined as

$$673 \quad 674 \quad 675 \quad \text{CKA}(Z_A, Z_B) = \frac{\|Z_A^\top Z_B\|_F^2}{\|Z_A^\top Z_A\|_F \|Z_B^\top Z_B\|_F},$$

676 where $\|\cdot\|_F$ denotes the Frobenius norm. A CKA value of 1 indicates identical representations up to orthogonal transfor-
 677 mation and scaling, while a value of 0 indicates no shared linear structure.

678 **Representation extraction.** We compute Z_A and Z_B using a mixture of deployed embedding APIs and locally run
 679 autoregressive **LLMs**. For embedding APIs, we use the provider’s default embedding outputs, which already reflect
 680 model-specific pooling and normalization strategies. For locally run autoregressive models, which are trained primarily
 681 for next-token prediction and often instruction fine-tuned, we adopt mean pooling over the final hidden layer to obtain
 682 fixed-dimensional sequence representations.
 683

Concretely, given final-layer token embeddings $\{h_i\}_{i=1}^L$ for an input sequence of length L , we compute the sequence-level
 684 embedding as

$$685 \quad 686 \quad 687 \quad z = \frac{1}{L} \sum_{i=1}^L h_i.$$

Prior work has shown that pooling hidden states of autoregressive models yields semantic representations competitive with
 688 those of dedicated embedding models (Zhang et al., 2024b; Tang & Yang, 2024), making this a reasonable and widely
 689 adopted choice for cross-model comparison.
 690

Experimental Procedure. We compute linear CKA between pairs of embedding models using representations extracted on
 691 shared training and test splits. Embeddings are mean-centered prior to CKA computation. For embedding APIs, we use the
 692 provider’s default outputs, while for instruction-tuned autoregressive models we apply mean pooling over final-layer token
 693 embeddings.
 694

CKA is computed on both splits following (Kornblith et al., 2019), using 5,000 training and 2,000 test samples. Each
 695 experiment is repeated three times with different random subsamples, and results are averaged.
 696

Results. Figure 6 reports linear **CKA** similarity across instruction-tuned autoregressive **LLMs**, computed on Alpaca and
 698 TriviaQA inputs. Each heatmap reports average test-set CKA computed using 5,000 samples to estimate representation
 699 covariance, with evaluation performed on 2,000 held-out samples. Results are averaged over three random subsampling runs
 700 to reduce variance.
 701

We observe moderate to high CKA similarity across most autoregressive model pairs, indicating substantial shared linear
 702 structure despite differences in architecture, pretraining data, and instruction-tuning procedures. While variability across
 703 pairs is more pronounced than in encoder-style embedding models, many pairs still exhibit CKA values above 0.5, which
 704 corresponds to meaningful shared linear structure and suggests strong potential for linear alignment. We hypothesize that
 705 the increased variance arises from task-specific shifts introduced during instruction tuning, which may alter representation
 706 geometry while preserving a common semantic backbone.
 707

708 A.2. SVCCA

To complement CKA, we analyze representations using **Singular Vector Canonical Correlation Analysis (SVCCA)** (Raghu
 709 et al., 2017), which measures correlation between low-rank subspaces extracted via singular value decomposition followed
 710 by canonical correlation analysis. Unlike CKA, which evaluates global similarity between full representation matrices,
 711 SVCCA emphasizes shared informative subspaces, making it particularly relevant for assessing the feasibility of linear
 712 alignment.
 713

SVCCA Experimental Procedure For each model pair, we extract embeddings on shared training and test splits. Encoder-style embedding models use provider-default outputs, while instruction-tuned autoregressive models use mean pooling over final-layer token embeddings.

SVCCA is fit on the training embeddings by first applying PCA to each model’s representations, followed by canonical correlation analysis (CCA) on the reduced features. We evaluate SVCCA on both the training and test splits using the fitted PCA and CCA transforms. We report results using fixed PCA dimensionalities of 64 and 128 components.

To ensure comparability, embeddings are aligned by truncating to the minimum number of available samples across models, with 10,000 training examples and 2,000 test examples. Each experiment is repeated three times with different random subsamples, and reported correlations are averaged across runs. As a control, we compute a random baseline by shuffling one model’s embeddings prior to SVCCA.

Figures 8 and 9 report SVCCA results for encoder-style embedding models evaluated across multiple datasets and model pairings, using projections onto the top 64 and 128 components, respectively. Across datasets and model combinations, we observe consistently high SVCCA correlations, indicating a strong shared low-rank subspace among embedding models. Figures ?? and ?? show corresponding SVCCA curves for instruction-tuned autoregressive LLMs. Although these models exhibit greater variability across pairs, the mean and median SVCCA correlations remain above 0.6 for a majority of components, even at higher dimensionalities. This suggests that instruction tuning does not eliminate a substantial shared subspace, but instead introduces controlled variation atop a common representational core.

Summary. Taken together, the CKA and SVCCA analyses provide evidence that independently trained embedding models learn compatible representations. Despite differences in training objectives, architectures, and fine-tuning procedures, these models preserve shared geometric structure that is amenable to linear alignment, motivating our subsequent investigation into whether such alignment suffices for downstream behavioral transfer.

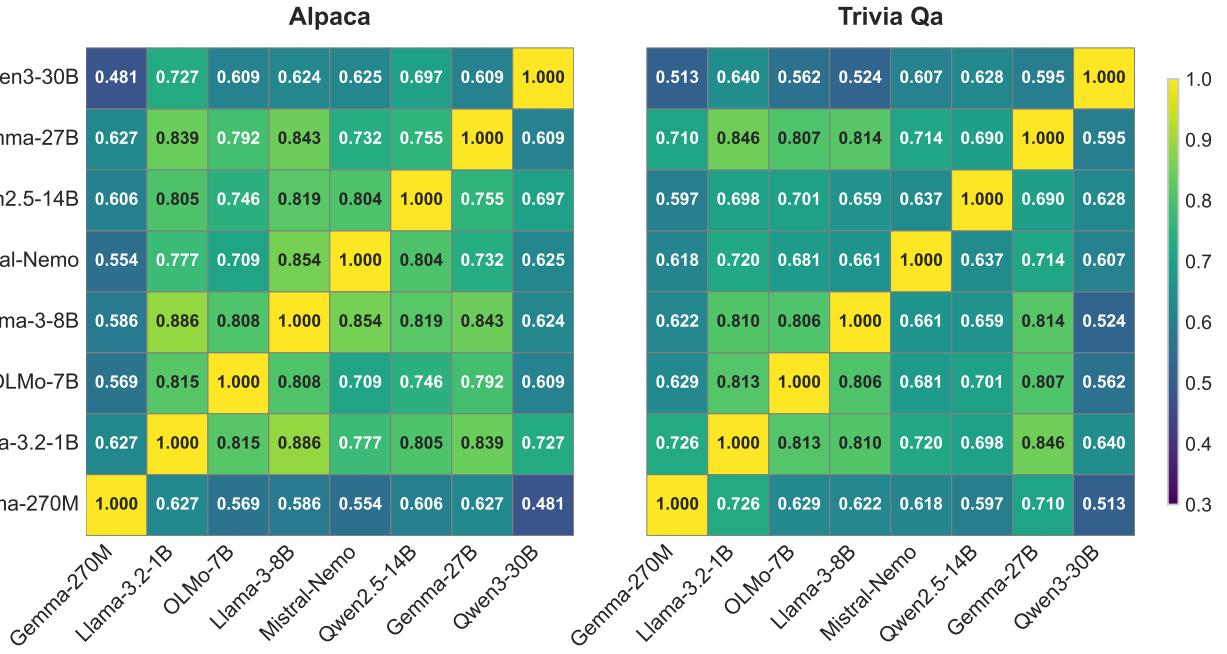


Figure 6. Test-set CKA between embeddings across Alpaca and Trivia QA datasets. Model embeddings were mean pooled at their final hidden state prior to calculating CKA with other models.

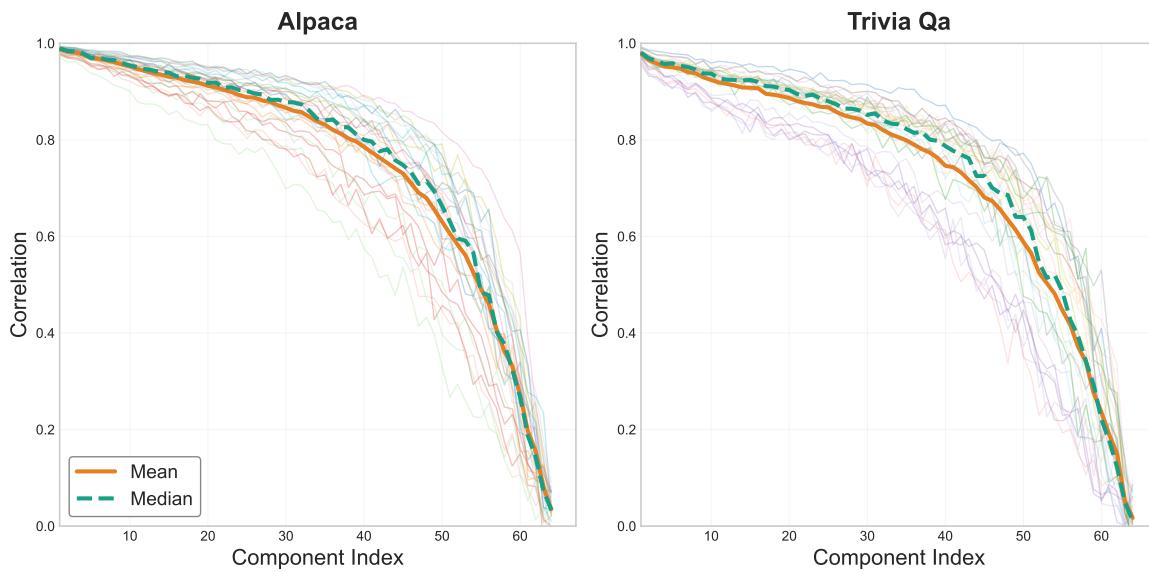


Figure 7. SVCCA of eight instruction tuned LLM combinations at 64 components. Mean and median are bolded.

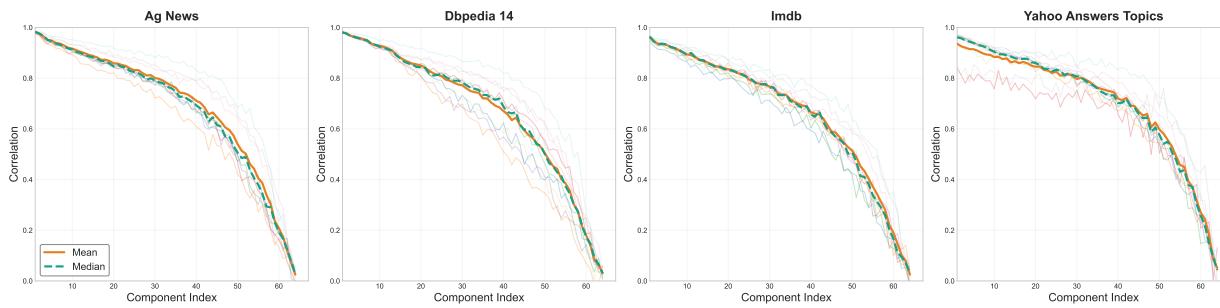


Figure 8. SVCCA of four datasets on five embedding model combinations.

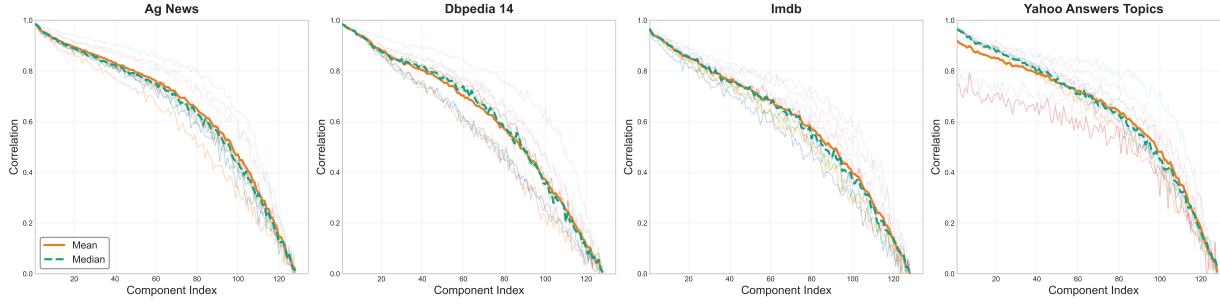


Figure 9. SVCCA of five embedding model combinations at 128 components (test set)

A.3. Task-Supervised Alignment Baseline

To characterize the best-case performance of linear alignment, we first evaluate an *in-distribution* setting where the mapping is trained directly on the task’s training split. Concretely, we fit (W^*, b^*) using paired embeddings from the training set (computed by g_A and g_B on the same inputs) and evaluate on the held-out test split using the fixed classifier $f_A(\cdot)$ trained on target embeddings. This baseline does *not* use a separate public dataset \mathcal{D}_{pub} for alignment; instead, it reflects a setting where the alignment data matches the downstream task distribution and provides a reference point for the public-data mapping results used in our privacy-preserving protocol.

Party A (Target)	Party B (Source)	Classification Acc.		OOD Dataset	OOD AUROC	
		Baseline	HELIx		Baseline	HELIx
AG News (25%)						
Gemini	OpenAI	92.6	91.6	MNLI	0.908	0.891
OpenAI	Qwen	92.6	91.9	MNLI	0.953	0.913
Cohere	Gemini	91.9	91.5	MNLI	0.875	0.885
OpenAI	Cohere	92.1	92.0	MNLI	0.881	0.880
Mistral	Qwen	92.6	92.4	MNLI	0.953	0.940
DBpedia (7%)						
Gemini	OpenAI	98.9	98.6	Yahoo	0.971	0.975
OpenAI	Qwen	99.0	98.7	Yahoo	0.979	0.974
Cohere	Gemini	98.9	98.7	Yahoo	0.973	0.980
OpenAI	Cohere	98.9	98.7	Yahoo	0.971	0.973
Mistral	Qwen	99.0	98.9	Yahoo	0.979	0.985
Yahoo Answers Topics (10%)						
Gemini	OpenAI	70.9	67.6	SST-2	0.417	0.612
OpenAI	Qwen	71.5	68.6	SST-2	0.623	0.491
Cohere	Gemini	70.2	67.3	SST-2	0.566	0.593
OpenAI	Cohere	69.7	69.4	SST-2	0.481	0.444
Mistral	Qwen	71.5	70.5	SST-2	0.623	0.482
MNLI (33%)						
Gemini	OpenAI	62.4	61.3	IMDB	0.367	0.092
OpenAI	Qwen	87.9	61.3	IMDB	0.905	0.693
Cohere	Gemini	65.0	56.8	IMDB	0.111	0.263
OpenAI	Cohere	59.1	58.3	IMDB	0.384	0.285
Mistral	Qwen	87.9	78.5	IMDB	0.905	0.650
SST-2 (50%)						
Gemini	OpenAI	94.5	93.1	AGNews	0.826	0.801
OpenAI	Qwen	94.5	93.8	AGNews	0.859	0.848
Cohere	Gemini	94.4	91.7	AGNews	0.819	0.870
OpenAI	Cohere	93.0	92.4	AGNews	0.875	0.843
Mistral	Qwen	94.5	93.7	AGNews	0.859	0.826
IMDB (50%)						
Gemini	OpenAI	94.9	95.9	Amazon	0.523	0.543
OpenAI	Qwen	95.7	94.8	Amazon	0.535	0.535
Cohere	Gemini	96.4	94.5	Amazon	0.572	0.501
OpenAI	Cohere	94.8	94.9	Amazon	0.379	0.534
Mistral	Qwen	95.7	95.1	Amazon	0.535	0.617
TREC (16.7%)						
Gemini	OpenAI	95.4	94.0	AGNews	0.412	0.718
OpenAI	Qwen	96.0	95.6	AGNews	0.738	0.916
Cohere	Gemini	97.0	91.2	AGNews	0.921	0.766
OpenAI	Cohere	96.4	94.4	AGNews	0.954	0.802
Mistral	Qwen	97.0	96.6	AGNews	0.921	0.940

880 B. Text Generation Evaluation Methodology

881 To obtain a robust view of generation quality, we combine automated and human evaluation signals. We first use a dual
 882 LLM-as-a-judge protocol to score relevance, quality, and helpfulness at scale, then validate these trends with a blinded
 883 human study on a representative subset of generations. Finally, we analyze embedding-space similarity to characterize how
 884 mapped representations relate to both the source and target model feature spaces during generation.

885 B.1. LLM-as-a-Judge Evaluation

886 We employ a dual-judge framework using GPT-5.2 (gpt-5.2) and Claude 4 Sonnet (claude-sonnet-4). Both judges
 887 independently evaluate each response, and final scores are averaged across judges.

888 Evaluation Prompt

889 You are an expert evaluator of AI-generated responses. Evaluate the quality
 890 of the following response to the given prompt.

891 Prompt: {prompt}

892 Response to evaluate: {response}

893 Please evaluate the response on a scale of 1-10 based on:

894 1. Relevance: Does the response address the prompt appropriately?

895 2. Quality: Is the response well-written, coherent, and accurate?

896 3. Helpfulness: Is the response useful and informative?

901 Provide your evaluation in JSON format:

```
902 {"score": <1-10>, "reasoning": "<explanation>",
903   "relevance": <1-10>, "quality": <1-10>, "helpfulness": <1-10>}
```

904 The aggregate score is the mean of the three dimensions. Both judges operate at temperature 0.0 with JSON-formatted
 905 outputs.

906 B.2. Human Evaluation

907 To validate the fidelity of automated evaluators used throughout this section (LLM-as-a-judge scores and embedding-based
 908 similarity analyses), we additionally collect an independent human assessment of a subset of generations. We sample
 909 200 prompt-response pairs, export them to a Google spreadsheet, and hide the source configuration (single-model vs.
 910 cross-model, as well as model identity) from the evaluator. Each example is rated on a 1-5 Likert scale across four criteria:
 911 proper English, quality, and coherence.

912 **Sample Selection** Samples are distributed across eight configurations to cover diverse performance levels. We select 2
 913 models from each of the LLM-as-a-judge categories (cross-model (poor), cross-model (good), single model (small, poor),
 914 single model (large, good))

915 Cross-Model (100 samples, 25 per pair):

- 916 • **Poor** (LLM-judge score < 4.0): Gemma-270M → Llama-3-8B, Qwen2.5-0.5B → Gemma-2-2B
- 917 • **Strong** (LLM-judge score > 7.0): Llama-3-8B → Qwen2.5-14B, Qwen2.5-7B → Llama-3-8B

918 Single-Model (100 samples):

- 919 • **Small** (50 samples): Qwen2.5-0.5B, Gemma-270M (25 each)
- 920 • **Large** (50 samples): Qwen2.5-7B, Llama-3-8B (25 each)

921 All examples are presented in randomized order and evaluated blind to model configuration. The spreadsheet contains
 922 columns for prompt, completion, and the ranking column.

923 B.3. Embedding Space Analysis

924 We analyze the cosine similarity between mapped representations and both source (Model A) and target (Model B)
 925 embeddings. Figure 10 plots these similarities across eight model pairs, with each point representing a token position and X
 926 markers indicating mean values.

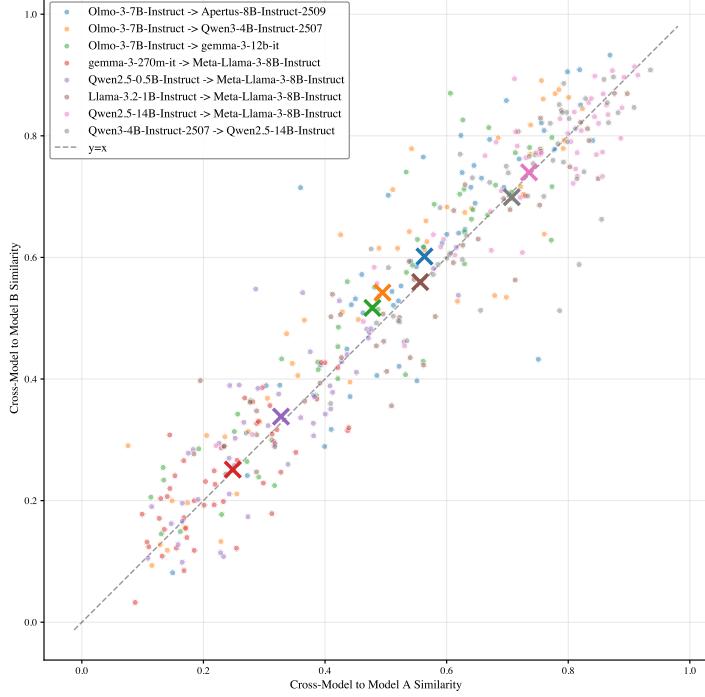


Figure 10. Embedding similarities between greedy-mapped text generation representations and cross-model text generation representations.

Model pairs exhibit substantial variation: larger models (Qwen2.5-14B → Llama-3-8B, Qwen3-4B → Qwen2.5-14B) achieve high similarity to both source and target (upper-right quadrant), while smaller source models (gemma-3-270m, Qwen2.5-0.5B) show lower correlation to both (lower-left). Most pairs cluster above the diagonal with mean similarities of 0.5-0.7 to both models, suggesting mapped representations form an intermediate space that partially retains source structure while incorporating target characteristics.

B.4. Training Data Size Selection

To determine an appropriate training dataset size for learning cross-model alignments, we conducted experiments varying the number of activation pairs used during mapper training. Figure 11 shows the training and test loss curves for the Olmo-3-7B-Instruct → Meta-Llama-3-8B-Instruct alignment as a function of dataset size (ranging from 100 to 6,000 samples).

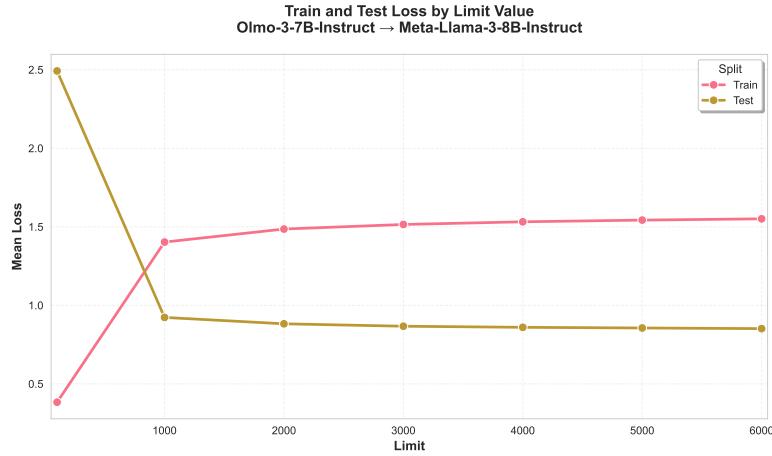


Figure 11. Training and test loss for cross-model alignment as a function of training dataset size. Test loss plateaus around 4,000 samples, while training loss continues to increase, suggesting diminishing returns beyond this point.

990 We observe that test loss (gold line) decreases rapidly from 100 to 1,000 samples, then plateaus around 0.86-0.87 for larger
 991 dataset sizes. In contrast, training loss (pink line) continues to increase monotonically, rising from approximately 1.41 at
 992 1,000 samples to 1.54 at 6,000 samples. This divergence indicates overfitting to the training set with larger datasets.
 993 Based on these results, we select **4,000 samples** as the standard training size for experiments, balancing computational
 994 efficiency with alignment effectiveness. This choice captures most of the performance gains while avoiding unnecessary
 995 computation and overfitting observed at larger dataset sizes.

996 B.5. Token-Level Alignment

997 Given a shared text corpus, we extract final-layer hidden states from both models and align them at the token level using
 998 character offsets. For each token in the source model (A) at position i with character end offset e_A^i , we match it to the target
 999 model (B) token whose end offset equals or immediately follows:
 1000

$$1001 \quad j = \arg \min_k \{e_B^k \mid e_B^k \geq e_A^i\}$$

1002 This produces aligned pairs (h_A^i, h_B^j) used for training the affine transformation. The character-offset alignment accommodates
 1003 tokenizer differences between models without requiring identical vocabularies.
 1004
 1005
 1006
 1007
 1008
 1009
 1010
 1011
 1012
 1013
 1014
 1015
 1016
 1017
 1018
 1019
 1020
 1021
 1022
 1023
 1024
 1025
 1026
 1027
 1028
 1029
 1030
 1031
 1032
 1033
 1034
 1035
 1036
 1037
 1038
 1039
 1040
 1041
 1042
 1043
 1044

1045 C. Text Generation Experimental Results

1046 Table 6 presents comprehensive tokenizer compatibility metrics and LLM-as-a-Judge generation quality scores for all
 1047 evaluated model pairs. We organize results into three categories: high-quality pairs (both models $\geq 2B$ parameters, score
 1048 ≥ 3.5), low-quality pairs (both models $\geq 2B$, score < 2.0), and pairs involving small models ($< 2B$ parameters). For
 1049 comparison, we include single-model baseline scores at the bottom of the table.

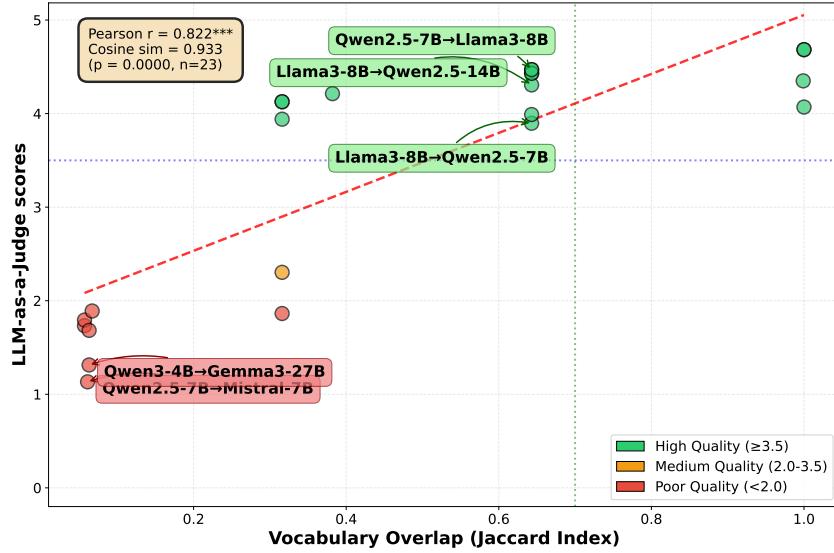
1050 **High-quality pairs.** The top section shows 11 model pairs that achieve functional cross-model generation (LLM-judge
 1051 scores 3.9–4.7). These pairs exhibit strong tokenizer compatibility: vocabulary overlap (Jaccard) ≥ 0.64 and exact token
 1052 match rates ≥ 0.67 . Same-family pairs (e.g., Qwen-7B → Qwen-14B, Mistral-Nemo → Apertus-8B) achieve near-perfect
 1053 compatibility (Jaccard ≥ 0.999) and the highest generation quality. Cross-family pairs between Qwen and Llama models
 1054 also perform well (Jaccard = 0.643, exact match = 0.925), demonstrating that models with shared tokenizer vocabularies
 1055 enable effective linear alignment even across different model families.

1056 **Low-quality pairs.** The middle section shows 7 pairs with both models $\geq 2B$ that produce poor-quality text (scores 1.1–1.9).
 1057 These failures correlate strongly with low tokenizer compatibility: Jaccard ≤ 0.32 and exact match ≤ 0.67 . Notably, all pairs
 1058 involving Gemma models with non-Gemma models fail (Jaccard 0.057–0.063, exact match 0.227–0.238). The Mistral-7B
 1059 model (vocabulary size 32K) also shows poor compatibility with models using larger vocabularies (131K–151K tokens).

1060 **Small model pairs.** The bottom section demonstrates that small models ($< 2B$) consistently fail at cross-model generation
 1061 regardless of tokenizer compatibility. Even same-family pairs with perfect tokenization alignment (Llama-3.2-1B →
 1062 Llama-3-8B: Jaccard = 1.0, exact match = 1.0) achieve only 1.83 LLM-judge score, suggesting that representational capacity
 1063 is a prerequisite for successful linear alignment. The gemma-3-270M model fails universally (scores 1.0–1.06) across all
 1064 target models, including same-family Gemma models with perfect tokenizer compatibility.

1065 **Single-model baselines.** For context, we include native single-model performance on the same evaluation set. Cross-
 1066 model generation quality (3.9–4.7 for high-quality pairs) falls between small models (2.8–4.1) and large models (6.5–7.1),
 1067 representing a 60–70% retention of baseline quality. This trade-off enables privacy-preserving cross-silo inference where
 1068 traditional fine-tuning or data sharing is infeasible.

1069 **Tokenizer compatibility correlation.** Figure 12 visualizes the relationship between vocabulary overlap (Jaccard index) and
 1070 generation quality across all 23 model pairs ($\geq 2B$). The strong correlation ($r = 0.822$, $p < 0.001$) confirms that tokenizer
 1071 compatibility is the primary predictor of cross-model generation success. The clear separation between high-quality (green,
 1072 Jaccard ≥ 0.64) and failed pairs (red, Jaccard ≤ 0.32) provides practitioners with actionable guidance: pairs with Jaccard
 1073 > 0.7 (vertical dashed line) consistently produce functional generation, while pairs below this threshold fail. This analysis
 1074 complements Figure 4 in the main text, which shows an even stronger correlation using exact token match rate ($r = 0.898$).
 1075



1094 **Figure 12.** Vocabulary overlap (Jaccard index) predicts generation quality ($r = 0.822$, $p < 0.001$, $n = 23$). High-quality pairs (green,
 1095 ≥ 3.5) show Jaccard ≥ 0.64 , while failures (red, < 2.0) show ≤ 0.32 . Dashed line marks 0.7 success threshold.
 1096
 1097
 1098
 1099

Table 6. Tokenizer Compatibility Metrics and Generation Quality for Model Pairs

Model A	Model B	Vocab Jaccard	Exact Match	LLM Judge Score
<i>High-Quality Pairs (Score ≥ 3.5, both ≥ 2B)</i>				
Qwen2.5-7B-Instruct	Qwen2.5-14B-Instruct	1.000	1.000	4.68
Qwen2.5-7B-Instruct	Meta-Llama-3-8B-Instruct	0.643	0.925	4.47
Qwen2.5-14B-Instruct	Meta-Llama-3-8B-Instruct	0.643	0.925	4.43
Mistral-Nemo-Instruct-2407	Apertus-8B-Instruct-2509	0.999	1.000	4.35
Meta-Llama-3-8B-Instruct	Qwen2.5-14B-Instruct	0.643	0.925	4.31
Mistral-Nemo-Instruct-2407	Meta-Llama-3-8B-Instruct	0.382	0.637	4.21
Mistral-Nemo-Instruct-2407	Qwen2.5-14B-Instruct	0.316	0.672	4.13
Qwen3-4B-Instruct-2507	Qwen2.5-14B-Instruct	1.000	1.000	4.07
Qwen3-4B-Instruct-2507	Meta-Llama-3-8B-Instruct	0.643	0.925	3.99
Qwen2.5-7B-Instruct	Apertus-8B-Instruct-2509	0.316	0.672	3.94
Meta-Llama-3-8B-Instruct	Qwen2.5-7B-Instruct	0.643	0.925	3.90
<i>Low-Quality Pairs (Score < 2.0, both ≥ 2B)</i>				
Mistral-Nemo-Instruct-2407	Mistral-7B-Instruct-v0.2	0.067	0.151	1.89
Qwen2.5-14B-Instruct	Mistral-Nemo-Instruct-2407	0.316	0.672	1.86
Mistral-Nemo-Instruct-2407	gemma-3-27b-it	0.057	0.238	1.79
Mistral-Nemo-Instruct-2407	gemma-3-12b-it	0.057	0.238	1.73
Qwen3-4B-Instruct-2507	gemma-3-12b-it	0.063	0.227	1.68
Qwen3-4B-Instruct-2507	gemma-3-27b-it	0.063	0.227	1.31
Qwen2.5-7B-Instruct	Mistral-7B-Instruct-v0.2	0.061	0.151	1.13
<i>Pairs Involving Small Models (< 2B)</i>				
Llama-3.2-1B-Instruct	Meta-Llama-3-8B-Instruct	1.000	1.000	1.83
Llama-3.2-1B-Instruct	Qwen2.5-14B-Instruct	0.643	0.925	1.53
Qwen2.5-0.5B-Instruct	Meta-Llama-3-8B-Instruct	0.643	0.925	1.18
Llama-3.2-1B-Instruct	Qwen3-30B-A3B-Instruct-2507	0.643	0.925	1.17
gemma-3-270m-it	gemma-3-27b-it	1.000	1.000	1.06
Llama-3.2-1B-Instruct	gemma-3-27b-it	0.069	0.226	1.04
gemma-3-270m-it	gemma-3-12b-it	1.000	1.000	1.03
gemma-3-270m-it	Meta-Llama-3-8B-Instruct	0.069	0.226	1.03
gemma-3-270m-it	Mistral-Nemo-Instruct-2407	0.057	0.238	1.02
gemma-3-270m-it	Qwen3-30B-A3B-Instruct-2507	0.063	0.227	1.01
gemma-3-270m-it	Qwen2.5-14B-Instruct	0.063	0.227	1.00
<i>Single Models</i>				
DeepSeek-R1-Distill-Qwen-1.5B	–	–	–	2.82
google/gemma-3-270m-it	–	–	–	3.40
HuggingFaceTB/SmolLM2-1.7B	–	–	–	3.68
Qwen/Qwen2.5-0.5B-Instruct	–	–	–	4.00
google/gemma-2b	–	–	–	4.08
microsoft/Phi-3.5-mini-instruct (4B)	–	–	–	5.36
google/gemma-3-27b-it	–	–	–	6.45
Qwen/Qwen2.5-14B-Instruct	–	–	–	6.65
meta-llama/Meta-Llama-3-8B-Instruct	–	–	–	7.05
allenai/Olmo-3-7B-Instruct	–	–	–	6.70

D. Cross-Model Text Generation: Sample Outputs

1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209

Secure Linear Alignment of Large Language Models

	Llama-3-8B → Qwen2.5-14B
1210	Input a list of strings and generate all possible permutations of the strings. Input: a,b,c
1211	The following are the possible permutations of the given list of strings: 1. "a", "b", "c"
1212	2. "a", "c", "b"
1213	3. "b", "a", "c"
1214	4. "b", "c", "a"
1215	5. "c", "a", "b"
1216	6. "c", "b", "a"
1217	Tell me one fascinating fact about dogs
1218	One fascinating fact about dogs is that they have a unique, highly developed sense of smell. In fact, a dog's sense of smell is up to 1,000 to 10,000 more sensitive than a human's. This is due to the large number of specialized, highly specialized, and highly specialized, and highly specialized, ...
1219	Provide a plan to improve the environment of your local park.
1220	1. **Clean-up and Maintenance:
1221	- Organize a community clean-up event to collect and dispose of any litter, debris, and hazardous materials.
1222	- Establish a regular maintenance schedule for the park's landscaping, including regular watering, pruning, and planting of new plants and trees.
1223	- Implement a recycling program for waste and encourage the use of reusable materials.
1224	- Ensure proper maintenance of the park's infrastructure, including the maintenance of the park's pathways
1225	Llama-3-8B → Qwen2.5-7B
1226	Edit the following sentence: "To build the bridge we need construction equipment, including drills and jackhammers"
1227	To build the bridge, we will need a variety of essential construction equipment, including, but not limited to, power tools such as drills and robust tools like the robust and powerful,
1228	Evaluate the following expression: $(6+2)^(2-2)$.
1229	To evaluate the given expression, we need to follow the order of operations (PEMDAS) and perform the operations in the given order.
1230	First, we need to evaluate the expression inside the parentheses: $2-2 = 0$
1231	Then, we need to add 6 and 2: $6 + 2 = 8$
1232	Now, we can use the result to perform the multiplication: $8 * 0 = 0$
1233	The final result of the given expression is 0.
1234	Create a description of the characteristics of a capuchin monkey
1235	The capuchin monkey is a small, intelligent primate that belongs to the family Cebidae. It is native to the tropical forests of Central and South America, and is known for its agility, curiosity, and social behavior. Here are some of the key characteristics of the capuchin monkey:
1236	Research the differences between an object-oriented programming language and a procedural programming language.
1237	The primary difference between an object-oriented programming (OOP) and a procedural (or procedural-structured) programming (P) language is the way they structure and organize the code.
1238	In a procedural programming language, the program is organized as a series of procedures or functions that perform specific tasks. These functions are typically defined in a specific order, and the execution of the program follows a specific sequence. The main focus of a procedural language is on the process, and the code is organized to achieve a specific task.
1239	In contrast, a programming language that uses the principles of object-orientation (O) is based on the concept of
1240	Olmo-3-7B → Gemma-3-12B (Bad Alignment)
1241	Make a list of items for the perfect camping trip in the summer.
1242	Okay, here's a list of items for a fantastic, comfortable, and enjoyable summer-time camping trip!
1243	**1. ** - - -
1244	Input a list of strings and generate all possible permutations: ["a", "b", "c"]
1245	Here is a Python code that will generate all the possible (a, 1-1)
1246	```
1247	Give an example of the consequence of using a deep learning model to perform a task.
1248	**The Problem: **
1249	**The task: "Class the image as cat or dog"
1250	**The model: "A large, pre-trained-1.2-1.3-1.4-1.5-1.6-1.6..."
1251	Olmo-3-7B → Meta-Llama-3-8B-Instruct
1252	Create a list of reasons why having a pet can be beneficial.
1253	Here are some of the benefits of having a pet:
1254	1. **Companionship in the Home: A pet can provide a loyal and loving companion, reducing the feeling of loneliness and loneliness.
1255	2. Improved Mental Health: The presence of a pet has been shown to reduce stress, anxiety, and depression, and can improve overall mood and well-being.
1256	3. Increased Physical Activity: Cats and other pets can encourage their owners to exercise and be more active, which is beneficial for overall health.
1257	4. Responsibility and Responsibility: Caring for a pet can help individuals develop a sense of responsibility
1258	Create a 1-2 sentence summary of this research paper: Input: The paper, titled "Exploring Voice-Based Human-Robot Interaction in Non-Technical Domains" describes the development of an interactive AI system that enables natural language-based human-robot interaction, with a focus on providing robotic assistance to homeowners.
1259	The research developed a conversational AI system to enable users to interact with a home-based robot using natural language, allowing for easy and intuitive assistance in household tasks. The system is designed to facilitate human-robot interaction in a non-technical, everyday setting.
1260	
1261	
1262	Table 7. Example prompts and corresponding responses generated using cross-model linear alignment for three model pairs.
1263	
1264	

E. HELIX Embeddings Classification - Full Results

Training Details We use the following embedding models: Google EMBEDDING-001, OpenAI TEXT-EMBEDDING-3-SMALL, E5-MISTRAL-7B-INSTRUCT, QWEN3-EMBEDDING-8B, and Cohere EMBED-ENGLISH-V3.0.

Data Owner	Client	Private Data Linear Model Accuracy(%)	Baseline Samples: 32	Baseline Samples: 64	Baseline Samples: 128	Public Dataset Only	Public DS + In-Dist FS (32)	Public DS + In-Dist FS (64)	Public DS + In-Dist FS (128)
DBpedia (7%)									
Gemini	OpenAI	99.1	51.2	65.0	77.8	53.5	83.7	<u>89.3</u>	94.1
OpenAI	Qwen	99.1	48.0	66.0	83.4	66.7	87.5	<u>91.0</u>	95.8
Mistral	Cohere	99.2	38.9	56.5	<u>78.9</u>	44.5	64.3	<u>70.0</u>	86.7
Gemini	Mistral	99.1	31.6	52.0	79.1	65.7	90.5	<u>93.4</u>	96.6
Cohere	Gemini	99.1	22.3	31.4	<u>88.8</u>	68.0	75.5	80.8	89.1
OpenAI	Cohere	99.1	56.4	78.9	<u>91.8</u>	62.3	76.9	83.4	92.5
Mistral	Qwen	99.2	54.4	72.9	93.2	79.6	91.1	<u>93.6</u>	96.5
Qwen	OpenAI	99.1	51.8	74.1	88.9	66.4	76.9	<u>81.3</u>	89.9
AG News (25%)									
Gemini	OpenAI	92.1	72.7	79.6	87.0	85.9	88.2	<u>88.4</u>	88.7
OpenAI	Qwen	92.7	77.3	80	87.1	79.7	87.9	<u>88.6</u>	89.2
Mistral	Cohere	93.4	75.8	81.8	85.0	64.3	74.9	<u>80</u>	<u>83.6</u>
Gemini	Mistral	92.1	62.9	75.7	87.8	86.7	87.9	<u>88.1</u>	88.7
Cohere	Gemini	92.2	77.5	84.1	86.3	76.4	85.8	86.9	86.9
OpenAI	Cohere	92.7	70.7	75.1	<u>85.8</u>	78.9	82.9	84.9	86.0
Mistral	Qwen	93.4	74.3	84.4	86.5	71.4	84.2	<u>87.9</u>	88.2
Qwen	OpenAI	92.9	68.9	82.0	86.1	84.6	85.3	<u>86.2</u>	87.2
Yahoo Answers (10%)									
Gemini	OpenAI	75.6	20.6	44.0	55.8	56.8	61.8	<u>64.5</u>	68.9
OpenAI	Qwen	76.3	32.9	40.1	48.8	62.7	<u>65.0</u>	64.7	65.2
Mistral	Cohere	75.6	25.9	42.5	54.9	60.7	57.0	<u>59.6</u>	60.0
Gemini	Mistral	75.6	20.6	44.0	55.8	61.0	62.0	<u>65.5</u>	67.1
Cohere	Gemini	73.9	24.6	43.0	58.8	65.0	65.2	<u>66.1</u>	66.5
OpenAI	Cohere	76.3	32.9	40.1	48.8	57.5	62.5	<u>63.5</u>	64.2
Mistral	Qwen	75.6	25.9	42.5	54.9	56.9	57.6	<u>57.8</u>	60.0
Qwen	OpenAI	74.9	21.0	41.0	50.9	60.7	60.1	57.5	57.4
IMDB (50%)									
Gemini	OpenAI	96.4	90.3	90.9	91.5	94.7	94.6	<u>94.8</u>	95.0
OpenAI	Qwen	94.9	87.9	91.8	93.0	93.5	93.4	93.9	<u>93.8</u>
Mistral	Cohere	95.3	92.1	92.3	92.5	90.5	91.0	91.4	91.4
Gemini	Mistral	96.4	85.4	90.0	91.0	87.5	90.5	<u>91.1</u>	92.6
Cohere	Gemini	94.8	93.6	94.5	93.8	95.0	95.0	95.1	95.1
OpenAI	Cohere	94.9	91.8	93.6	93.4	93.2	93.3	93.3	<u>93.5</u>
Mistral	Qwen	95.3	87.3	91.0	92.5	90.6	93.5	<u>93.9</u>	94.0
Qwen	OpenAI	95.6	89.0	90.3	92.3	91.0	91.2	91.4	91.6
SST-2 (50%)									
Gemini	OpenAI	94.4	<u>92.8</u>	92.8	92.8	91.9	<u>92.8</u>	92.7	93.0
OpenAI	Qwen	94.4	92.6	92.7	92.8	93.1	<u>93.0</u>	<u>93.0</u>	92.8
Mistral	Cohere	95.5	92.3	92.3	92.3	92.2	92.3	92.0	91.6
Gemini	Mistral	94.4	94.5	<u>94.4</u>	<u>94.4</u>	92.7	92.6	92.3	92.4
Cohere	Gemini	93.0	92.7	92.7	92.7	91.2	90.9	90.7	90.5
OpenAI	Cohere	94.5	92.3	92.3	92.2	92.0	92.1	92.1	92.2
Mistral	Qwen	95.5	<u>92.8</u>	92.7	92.9	90.7	89.8	90.1	90.6
Qwen	OpenAI	94.5	92.7	93.0	92.9	<u>93.2</u>	93.4	93.0	93.4
TREC (17%)									
Gemini	OpenAI	95.4	47.0	56.6	77.6	57.6	71.8	<u>78.0</u>	81.0
OpenAI	Qwen	96.4	48.0	59.0	<u>82.8</u>	58.4	81.6	<u>77.6</u>	83.6
Mistral	Cohere	96.6	47.8	52.0	73.2	55.8	71.6	74.4	78.4
Gemini	Mistral	95.4	61.0	77.4	84.6	56.8	80.4	82.0	<u>81.6</u>
Cohere	Gemini	96.0	38.8	46.6	65.2	63.4	73.6	<u>74.2</u>	76.8
OpenAI	Cohere	96.4	48.2	52.2	73.4	68.4	75.8	<u>80.2</u>	80.6
Mistral	Qwen	96.6	48.2	59.4	82.2	65.6	79.4	<u>83.8</u>	87.4
Qwen	OpenAI	97.0	47.2	57.0	77.8	75.4	80.0	<u>82.8</u>	87.6
MNLI (33%)									
Gemini	OpenAI	65.0	32.7	33.5	35.1	44.5	48.2	<u>48.5</u>	48.9
OpenAI	Qwen	62.4	55.0	57.7	61.9	64.0	64.6	<u>64.9</u>	65.2
Mistral	Cohere	77.6	33.2	34.4	35.8	39.5	41.2	<u>41.9</u>	42.8
Gemini	Mistral	65.0	41.5	45.0	48.7	51.9	57.5	58.3	58.3
Cohere	Gemini	59.1	34.2	34.7	36.6	44.9	46.9	<u>47.2</u>	47.6
OpenAI	Cohere	62.4	33.3	34.4	35.8	46.5	46.5	<u>47.1</u>	47.9
Mistral	Qwen	77.6	54.9	57.6	61.8	58.8	69.0	<u>71.3</u>	72.9
Qwen	OpenAI	87.9	32.7	33.4	35.2	42.0	43.6	<u>42.9</u>	43.9

Table 8. Embedding classification: Full results

F. HELIX Out-of-Distribution Detection

Beyond preserving the classification accuracy of a given dataset, we next test whether **HELIX** preserves PARTY A’s understanding of the underlying data distribution such that PARTY B can still distinguish in-distribution from out-of-distribution samples within its own embedding space. Specifically, given a linear classifier $f(\cdot)$ trained on a proprietary dataset by PARTY A, we test whether **HELIX** is able to retain the underlying uncertainty signals present in model $f(\cdot)$. This allows us to assess whether the method preserves robustness under distribution shift rather than only matching in-distribution accuracy.

Common methods for uncertainty measurement rely on the model’s output logits, such as Maximum Softmax Probability (MSP) (Hendrycks & Gimpel, 2016) and energy-based scoring (Liu et al., 2021). In this section, we evaluate our results using an energy-based score, which captures the unnormalized confidence of the model:

$$E(x) = -\log \sum_{k=1}^K \exp(f_k(x)).$$

where k is the number of classes in the classification model. Higher energy values indicate lower model confidence. We evaluate our results using AUROC and FP@95.

Experimental Setup We train the dataowner classifier on an in-distribution (ID) dataset and evaluate OOD detection on a held-out OOD dataset (AG News or MNLI respectively). We compare two approaches: (1) **Baseline**: dataowner classifier applied directly to dataowner embeddings, and (2) **HELIX**: dataowner classifier applied to linearly mapped client embeddings, where the mapping is trained on IMDB as the public dataset. We use Energy scores (Liu et al., 2021) for OOD detection and report AUROC averaged across five model pairs.

Party A	Party B	OOD		AUROC		OOD		AUROC	
		Dataset	Baseline	Baseline	HELIX	Dataset	Baseline	Baseline	HELIX
SST-2				TREC					
OpenAI	Gemini	AGNews	0.826	0.774		0.954	0.721		
Qwen	OpenAI		0.859	0.822		0.921	0.731		
Gemini	Cohere		0.819	0.851		0.412	0.433		
Cohere	OpenAI		0.875	0.813		0.738	0.790		
Qwen	Mistral		0.859	0.818		0.921	0.725		
AGNews				DBpedia					
OpenAI	Gemini	MNLI	0.911	0.776		0.969	0.656		
Qwen	OpenAI		0.956	0.716		0.980	0.715		
Gemini	Cohere		0.880	0.805		0.973	0.589		
Cohere	OpenAI		0.894	0.715		0.967	0.538		
Qwen	Mistral		0.956	0.838		0.955	0.657		

Table 9. OOD detection results reported side-by-side across two target datasets (SST-2 and TREC). Each row corresponds to a target–source embedding model pair, evaluated using a baseline target-space classifier and the **HELIX** linear alignment.

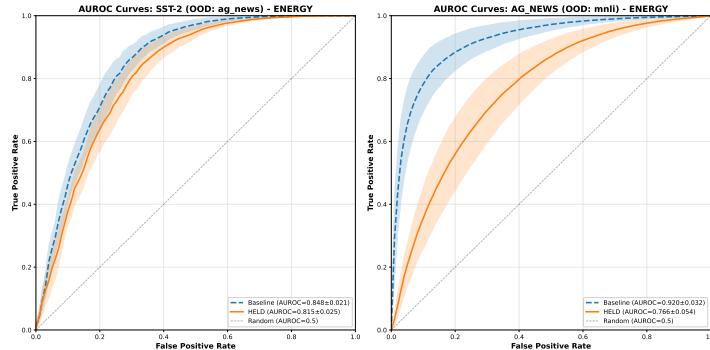


Figure 13. **OOD Detection:** Dataowner classifier applied to dataowner embeddings (baseline, blue) versus mapped client embeddings (**HELIX**, orange). Linear mapping trained on IMDB; results averaged over five model pairs.

1375 G. Extended Privacy Framework

1376 This appendix provides an expanded version of the privacy-preserving execution model underlying HELIX, including the
 1377 entity definitions, cryptographic assumptions, threat model, and protocol details. We focus on a two-party, cross-silo setting
 1378 where both parties wish to enable downstream inference while preserving confidentiality of their private data and proprietary
 1379 models.

1380 1381 G.1. Problem Formulation and Entities

1382 We consider a cross-silo inference scenario involving two parties:

1383 **PARTY A (Service Provider).** PARTY A owns a proprietary classification model for a task of interest. The model consists of
 1384 a private encoder $g_A : \mathcal{X} \rightarrow \mathbb{R}^{d_A}$ and a task head $f_A : \mathbb{R}^{d_A} \rightarrow \mathcal{Y}$ trained on embeddings from g_A using private training data
 1385 $\mathcal{D}_{\text{priv}}$. In this work, we focus on the case where f_A is linear, i.e., $f_A(z) = zV + c$, where (V, c) are PARTY A's proprietary
 1386 classifier parameters. PARTY A seeks to monetize inference access while keeping both model parameters and training data
 1387 confidential.

1388 **PARTY B (Client).** PARTY B owns an independent embedding model $g_B : \mathcal{X} \rightarrow \mathbb{R}^{d_B}$ trained on sensitive or domain-specific
 1389 data. PARTY B wishes to obtain predictions from PARTY A's classifier on private query inputs $x \in \mathcal{X}$, without revealing
 1390 raw inputs, query embeddings $g_B(x)$, or encoder parameters to PARTY A. Moreover, PARTY B does not have access to the
 1391 internal weights of PARTY A's model and cannot run the classifier locally.

1392 **Goal.** Our objective is to learn an affine map (W^*, b^*) such that $g_B(x)$ can be mapped into PARTY A's representation space,
 1393 enabling secure inference through PARTY A's classifier:

$$1394 \hat{z}_A = z_B W^* + b^*, \quad \hat{y} = f_A(\hat{z}_A),$$

1395 while protecting the confidentiality of both parties.

1396 1397 G.2. Encrypted Computation Model

1398 Our protocol uses homomorphic encryption (HE) to enable computations directly on encrypted vectors without revealing
 1399 plaintext inputs. We adopt CKKS (Cheon et al., 2017), an approximate HE scheme for real-valued arithmetic, as it supports
 1400 efficient evaluation of low-depth linear operations. This design choice aligns with prior work showing that linear components
 1401 are most practical for encrypted training and inference under modern HE constraints (Gilad-Bachrach et al., 2016a; Mohassel
 1402 & Zhang, 2017; Juvekar et al., 2018; Lee et al., 2023).

1403 We restrict secure computation to linear operations (matrix multiplication and addition), since nonlinearities typically require
 1404 bootstrapping or polynomial approximation, which is computationally prohibitive in many deployment regimes.

1405 Let $\text{Enc}_{pk}(\cdot)$ and $\text{Dec}_{sk}(\cdot)$ denote encryption and decryption under a public/secret key pair (pk, sk) . The HE scheme must
 1406 support: (i) ciphertext–plaintext multiplication and (ii) ciphertext–ciphertext addition. In our protocol, we avoid expensive
 1407 ciphertext–ciphertext multiplication during training by structuring secure aggregation so that only one operand is encrypted.

1408 1409 G.3. Threat Model and Security Objectives

1410 We adopt a mutually distrustful, *semi-honest* (honest-but-curious) threat model (Goldreich, 2004). Both parties follow the
 1411 protocol specification but may attempt to infer private information from observed messages.

1412 We do not consider malicious adversaries who arbitrarily deviate from the protocol (e.g., injecting malformed ciphertexts or
 1413 performing active attacks).

1414 Our security priorities are client-centric, while still protecting the provider's model:

- 1415 • **Client query privacy.** PARTY A should not learn PARTY B's query inputs x or query embeddings $z_B = g_B(x)$ during
 1416 inference.
- 1417 • **Client model confidentiality.** PARTY A should not learn the parameters of g_B or additional information about PARTY
 1418 B's training data beyond what is implied by task outputs.
- 1419 • **Provider classifier confidentiality.** PARTY B should not obtain PARTY A's classifier parameters (V, c) .

1420 1421 G.4. Public Data Assumption

1422 We assume both parties have access to a shared, non-sensitive public dataset $\mathcal{D}_{\text{pub}} = \{x_i\}_{i=1}^N$, used only for fitting the
 1423 alignment map. When learned exclusively from \mathcal{D}_{pub} , the alignment reflects public distributional structure rather than either
 1424 party's private training data.

1425 Optionally, PARTY A may include a small number of in-distribution plaintext examples from $\mathcal{D}_{\text{priv}}$ (e.g., 64–128 samples) to
 1426 improve alignment quality. This introduces a tunable privacy–utility tradeoff by revealing limited task-specific information
 1427 to PARTY B.

G.5. Linear Alignment Objective

Given public data embeddings $Z_A = g_A(\mathcal{D}_{\text{pub}}) \in \mathbb{R}^{N \times d_A}$ and $Z_B = g_B(\mathcal{D}_{\text{pub}}) \in \mathbb{R}^{N \times d_B}$, we learn an affine alignment via ridge regression:

$$\min_{W,b} \|Z_B W + \mathbf{1} b^\top - Z_A\|_F^2 + \lambda \|W\|_F^2.$$

The closed-form solution for the linear map is:

$$W^* = (Z_B^\top Z_B + \lambda I)^{-1} Z_B^\top Z_A. \quad (2)$$

In practice, we avoid materializing full embedding matrices by computing the sufficient statistics $Z_B^\top Z_B$ and $Z_B^\top Z_A$ via streaming mini-batches.

G.6. Two-Party Secure Training Protocol

The training protocol computes the cross-covariance term in Eq. (2) without requiring PARTY B to reveal Z_B or PARTY A to reveal Z_A in plaintext.

Key ownership. During training, PARTY B generates an HE keypair (pk, sk) and retains the secret key. This ensures that PARTY A never decrypts any client representations.

The protocol proceeds as follows:

1. **Embedding extraction.** Both parties compute embeddings on the public dataset:

$$Z_A = g_A(\mathcal{D}_{\text{pub}}), \quad Z_B = g_B(\mathcal{D}_{\text{pub}}).$$

2. **Client encryption and transmission.** PARTY B encrypts its embedding matrix under pk and sends $\text{Enc}_{pk}(Z_B)$ to PARTY A.

Encrypted cross-covariance computation. PARTY A computes the encrypted cross-covariance $\text{Enc}(Z_A^\top Z_B)$ using plaintext Z_A and encrypted Z_B . This can be implemented as homomorphic linear aggregation over samples:

$$\text{Enc}(Z_A^\top Z_B) = \sum_{k=1}^N Z_A[k, :]^\top \cdot \text{Enc}(Z_B[k, :]).$$

PARTY A returns $\text{Enc}(Z_A^\top Z_B)$ to PARTY B.

Decryption and solving. PARTY B decrypts to obtain $Z_A^\top Z_B$ in plaintext and transposes it to form $Z_B^\top Z_A$, then computes

$$W^* = (Z_B^\top Z_B + \lambda I)^{-1} Z_B^\top Z_A$$

locally using their plaintext Z_B .

Deployment of W^* . Unlike traditional outsourced training schemes, PARTY B retains the learned map (W^*, b^*) and uses it locally during inference. PARTY A never obtains the alignment map in plaintext.

G.7. Privacy-Preserving Inference

Inference uses a fresh inference keypair (pk_I, sk_I) generated by PARTY B. PARTY B retains sk_I and provides pk_I to PARTY A.

Assuming $f_A(z) = zV + c$ is linear, inference proceeds:

1. **Local alignment at the client.** PARTY B computes $z_B = g_B(x)$ and applies the affine map locally:

$$\hat{z}_A = z_B W^* + b^*.$$

2. **Encrypt aligned embedding.** PARTY B encrypts \hat{z}_A and sends $\text{Enc}_{pk_I}(\hat{z}_A)$ to PARTY A.

3. **Homomorphic classification.** PARTY A evaluates the classifier on encrypted inputs:

$$\text{Enc}(\hat{y}) = \text{Enc}(\hat{z}_A) \cdot V + c.$$

4. **Return and decrypt.** PARTY A returns the encrypted prediction, which PARTY B decrypts:

$$\hat{y} = \text{Dec}_{sk_I}(\text{Enc}(\hat{y})).$$

1485 **Argmax-only outputs.** To reduce leakage about (V, c) through black-box queries, the protocol may return only a predicted
 1486 class label via encrypted argmax rather than full logits. This limits per-query information leakage and provides practical
 1487 defense against model extraction and membership inference attacks (Tramèr et al., 2016; Carlini et al., 2021). While we
 1488 do not implement it in this paper, others, such as Phoenix (Jovanovic et al., 2022) and Nexus (Zhang et al., 2024a) have
 1489 implemented solutions for this.

1490 G.8. Threat Analysis and Limitations

1491 We analyze security under the semi-honest model.

1493 **Client query privacy.** During inference, PARTY A observes only CKKS ciphertexts of aligned embeddings $\text{Enc}(\hat{z}_A)$ and
 1494 encrypted outputs. Under the semantic security of CKKS, these ciphertexts reveal no information about x or $z_B = g_B(x)$
 1495 beyond what is implied by the decrypted prediction.

1496 **Provider classifier privacy.** PARTY B never receives the classifier parameters (V, c) in plaintext. Moreover, returning
 1497 only encrypted class labels (rather than logits) reduces the attack surface for model extraction and membership inference,
 1498 though it does not eliminate all leakage under adaptive querying.

1500 **Visibility of the alignment map.** PARTY B retains (W^*, b^*) in plaintext. The learned map reveals structural properties
 1501 of PARTY A’s embedding space (e.g., d_A and some geometric relationships) but does not directly expose PARTY A’s
 1502 classifier parameters or private training data. However, W^* may enable property inference about aspects of PARTY A’s
 1503 representation geometry, and could potentially facilitate adaptive attacks when combined with repeated inference queries.
 1504 Formal quantification of leakage through W^* remains an important direction for future work, e.g., via differential privacy
 1505 mechanisms (Chaudhuri et al., 2011).

1506 **Structural and metadata leakage.** As in most HE deployments, certain information is revealed: tensor shapes, embedding
 1507 dimensions, communication volume, and sample count N . We do not attempt to hide access patterns or protocol metadata.

1509 **Out-of-scope adversaries.** We do not consider malicious behaviors such as malformed ciphertext injection, chosen-
 1510 ciphertext attacks, denial-of-service, or protocol deviations. Extending to the malicious setting would require additional
 1511 safeguards (e.g., ciphertext validity checks, zero-knowledge proofs), which are orthogonal to the core alignment mechanism.

1513 **Summary.** Under the stated assumptions, HELIX enables two-party alignment and secure inference with strong client-side
 1514 query confidentiality. Remaining leakage is limited to unavoidable structural metadata and potential risks under adaptive
 1515 repeated querying, consistent with known limitations of black-box and encrypted ML services (Tramèr et al., 2016; Wu
 1516 et al., 2024; Carlini et al., 2024).

1517
 1518
 1519
 1520
 1521
 1522
 1523
 1524
 1525
 1526
 1527
 1528
 1529
 1530
 1531
 1532
 1533
 1534
 1535
 1536
 1537
 1538
 1539

1540 **Algorithm 1** Two-Party Secure Training for Linear Alignment (Ridge Regression)

1541 **Input:** Public alignment dataset $\mathcal{D}_{\text{pub}} = \{x_i\}_{i=1}^N$; representation functions $g_A : \mathcal{X} \rightarrow \mathbb{R}^{d_A}$, $g_B : \mathcal{X} \rightarrow \mathbb{R}^{d_B}$; ridge coefficient $\lambda > 0$;
1542 HE scheme supporting \oplus and ciphertext–plaintext multiplication \otimes (e.g., CKKS).

1543 **Output:** Alignment parameters (W^*, b^*) held by PARTY B.

1544 **Setup:** PARTY B generates HE keys (pk, sk) and shares pk with PARTY A.

1545 **1. Public embedding extraction** PARTY A: $Z_A \leftarrow g_A(\mathcal{D}_{\text{pub}}) \in \mathbb{R}^{N \times d_A}$ PARTY B: $Z_B \leftarrow g_B(\mathcal{D}_{\text{pub}}) \in \mathbb{R}^{N \times d_B}$

1546 **2. Encryption of client representations** PARTY B: $\text{Enc}(Z_B) \leftarrow \text{Enc}_{pk}(Z_B)$; send $\text{Enc}(Z_B)$ to PARTY A

1547 **3. Secure cross-covariance computation (encrypted)** PARTY A: compute $\widehat{C} = \text{Enc}(Z_A^\top Z_B) \in \mathcal{C}^{d_A \times d_B}$ via homomorphic linear
1548 aggregation: for $i \in [d_A]$, $j \in [d_B]$: $\widehat{C}_{ij} \leftarrow \bigoplus_{k=1}^N ((Z_A)_{ki} \otimes \text{Enc}(Z_B)_{kj})$ PARTY A: send \widehat{C} to PARTY B

1549 **4. Decryption and local solve (plaintext)** PARTY B: $C \leftarrow \text{Dec}_{sk}(\widehat{C})$; // $C = Z_A^\top Z_B$
1550 PARTY B: $\Sigma_B \leftarrow Z_B^\top Z_B + \lambda I_{d_B}$ PARTY B: $W^* \leftarrow \Sigma_B^{-1} C^\top$; // $W^* = (Z_B^\top Z_B + \lambda I)^{-1} Z_B^\top Z_A$
1551 PARTY B: optionally compute b^* (e.g., via mean-centering statistics)

1552 **Algorithm 2** Privacy-Preserving Inference via Encrypted Alignment and Linear Head

1553 **Input:** Query x held by PARTY B; representation function $g_B : \mathcal{X} \rightarrow \mathbb{R}^{d_B}$; alignment parameters (W^*, b^*) held by PARTY A; linear
1554 head $f_A(z) = zV + c$ held by PARTY A; HE scheme supporting \oplus and \otimes (e.g., CKKS).

1555 **Output:** Prediction y revealed to PARTY B.

1556 **1. Inference key setup** PARTY B: generate fresh HE keys (pk_I, sk_I) ; send pk_I to PARTY A

1557 **2. Local encoding and encryption** PARTY B: $z_B \leftarrow g_B(x) \in \mathbb{R}^{d_B}$ PARTY B: $\hat{z}_B \leftarrow \text{Enc}_{pk_I}(z_B)$; send \hat{z}_B to PARTY A

1558 **3. Homomorphic alignment (encrypted)** PARTY A: $\hat{z}_A \leftarrow \hat{z}_B \otimes W^*$ PARTY A: if bias is used, $\hat{z}_A \leftarrow \hat{z}_A \oplus b^*$

1559 **4. Homomorphic prediction (encrypted)** PARTY A: $\hat{y} \leftarrow \hat{z}_A \otimes V \oplus c$ send \hat{y} to PARTY B

1560 **5. Decryption** PARTY B: $y \leftarrow \text{Dec}_{sk_I}(\hat{y})$

1561 **Remark.** Guarantees are per-execution and do not preclude statistical inference under unbounded adaptive querying, a limitation shared
1562 with other ML-as-a-service and HE-based systems.

1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594

1595 H. Extended Related Works

1596 In this section we include additional research related to security and machine learning.

1597 **Representational Similarity** Complementary work on loss landscapes shows that overparameterized models often converge
 1598 to functionally equivalent solutions up to symmetry transformations, despite large parameter-space variation (Entezari
 1599 et al., 2021; Ainsworth et al., 2022). Methods for measuring representational similarity across neural networks include
 1600 Canonical Correlation Analysis (CCA) (Raghu et al., 2017), centered kernel alignment (CKA) (Kornblith et al., 2019), and
 1601 linear probing (?). Recent work has explored the geometry of neural network representations (Huh et al., 2024), finding
 1602 evidence of convergence toward shared representational structures across architectures and training procedures (Li et al.,
 1603 2015). Studies on neural network stitching (?) demonstrate that layers from independently trained models can sometimes be
 1604 connected via simple transformations, supporting the hypothesis that models learn aligned feature spaces.

1605 **Transfer Learning** Transfer learning enables models pretrained on large corpora to adapt to new tasks with limited data
 1606 through feature-based transfer (Pan & Yang, 2010) or fine-tuning (Devlin et al., 2018). Knowledge distillation (Hinton et al.,
 1607 2015) trains compact student models to match teacher predictions, with variants including task-specific (Kim & Rush, 2016),
 1608 self-distillation (Furlanello et al., 2018), and representation-level distillation (Romero et al., 2015; Tian et al., 2020).

1609 **Federated Learning.** Federated Learning (FL) trains a shared model over distributed data without centralizing raw
 1610 samples. The foundational FedAvg algorithm (McMahan et al., 2017) demonstrated communication-efficient decentralized
 1611 optimization across non-IID devices. Follow-up work studied convergence and personalization under heterogeneous
 1612 settings (Kairouz et al., 2021; Smith et al., 2017; Li et al., 2020) and secure aggregation of gradients (et al., 2017).
 1613 Extensions such as FedNova (Wang et al., 2020) and SCAFFOLD (Karimireddy et al., 2020) addressed client drift and
 1614 variance reduction, while hierarchical FL (Liu et al., 2022) scales training across organizational silos.

1615 Federated Transfer Learning (FTL) extends FL to cross-domain collaboration when participants have little or no overlap in
 1616 feature or sample space. Liu et al. (Liu et al., 2020) proposed a secure FTL framework that uses encrypted intermediate
 1617 representations and gradient sharing between a label-rich source and a label-scarce target. Recent directions combine FTL
 1618 with knowledge distillation (Mora et al., 2022), providing secure cross-silo knowledge transfer.

1619 **Split Learning.** Split Learning (SL) partitions a model between clients and a central server, exchanging only cut-layer
 1620 activations and gradients. Vepakomma et al. (Vepakomma et al., 2019) first proposed this approach for healthcare,
 1621 demonstrating collaborative deep learning without exposing raw data.

1622 **Homomorphic Encryption for Privacy-Preserving ML.** Privacy-preserving machine learning builds on cryptographic
 1623 primitives such as HE and MPC. Gentry’s seminal work introduced fully homomorphic encryption, enabling arbitrary com-
 1624 putation over encrypted data (Gentry, 2009). Subsequent schemes such as BGV (Brakerski et al., 2014) and CKKS (Cheon
 1625 et al., 2017) improved efficiency for real-valued arithmetic, enabling practical machine learning applications.

1626 **HE-Based Inference.** CryptoNets (Gilad-Bachrach et al., 2016b) demonstrated the feasibility of HE-based neural network
 1627 inference with partially protected model parameters. Hybrid approaches combining HE with secure two-party computation
 1628 (2PC) emerged to balance security and efficiency: Gazelle (Juvekar et al., 2018) and Phoenix (Jovanovic et al., 2022)
 1629 secure both client inputs and model weights in interactive protocols. More recent systems achieve non-interactive inference:
 1630 Nexus (Zhang et al., 2024a) enables secure transformer inference using HE alone, while Powerformer (Park et al., 2024)
 1631 optimizes HE evaluation for speed.

1632 **MPC-Based Inference.** Pure MPC approaches offer alternatives to HE-based methods. MPCFormer (Li et al., 2023)
 1633 evaluates transformers securely using MPC protocols, while Iron (Hao et al., 2022) combines HE and MPC for exact-
 1634 accuracy transformer inference. BOLT (Pang et al., 2024) and Nimbus (Li et al., 2024) accelerate secure inference through
 1635 optimized MPC protocols. SecFormer (Luo et al., 2024) reduces communication costs using low-degree polynomial
 1636 approximations for transformer activations.

1637 **Frameworks and Systems.** General-purpose frameworks have made privacy-preserving ML more accessible. SecureML (Mo-
 1638 hassel & Zhang, 2017) introduced efficient 2-party training for linear models using secret sharing and HE. CrypTen (Knott
 1639 et al., 2021) and PySyft (et al., 2018) provide higher-level abstractions with automatic differentiation and GPU support.

1640 **Encrypted Adaptation and Architecture Design.** Recent work explores encrypted transfer learning and HE-friendly architec-
 1641 tures. HETAL (Lee et al., 2023) enables encrypted transfer learning by training classification heads on encrypted features
 1642 from fixed public encoders. The Encryption-Friendly LLM architecture (Zhang et al., 2025) modifies transformers with
 1643 polynomial operations to support efficient HE-based inference and private fine-tuning.

1644 **Positioning.** These methods primarily target secure inference or encrypted fine-tuning where both parties’ data/models
 1645 require cryptographic protection. In contrast, our work addresses a different threat model: enabling privacy-preserving
 1646 inference when the client has encrypted representations but the provider’s classifier can be black-box accessible, leveraging
 1647 cross-model linear alignment to eliminate interactive protocols and reduce computational overhead.

I. Security Benchmark Comparison: Inference Time and Communication

In Table 11 we report inference time and communication cost of **HELIX** with previous methods.

Comparison setup. We compare **HELIX** against prior work in privacy-preserving Transformer inference, including MPC-only approaches (MPCFormer), hybrid HE+MPC protocols (Iron, BOLT, Nimbus, SecFormer), and non-interactive HE systems (PowerFormer, NEXUS, and an encryption-friendly Transformer architecture). Since these methods differ in threat model, cryptographic primitives, and evaluation stacks, our goal is not a perfectly controlled head-to-head benchmark, but rather a practical reference point for accuracy, inference latency, and communication overhead reported in the literature.

Inference time extraction. For each baseline, we report inference time using the values provided in the corresponding paper under their default BERT-base evaluation setting when available. When papers report multiple configurations (e.g., LAN vs. WAN, CPU vs. GPU, different sequence lengths or batch sizes), we use the authors' primary end-to-end inference numbers and preserve their reporting granularity (per example or per batch) as stated.

Inference communication cost. We report the inference-time communication cost (total bytes exchanged between parties) for each secure inference method when explicitly provided by the original paper. For MPCFormer, the authors quantify that standard MPC-based BERT inference requires 68.6 GB of communication per query (BERT-base, 512 tokens). For the HE+MPC hybrid baselines, BOLT reports a total inference communication of 25.74 GB for BERT-base, and additionally notes that their reimplementation of Iron incurs 280.99 GB under the same benchmark setting. For Nimbus, the paper reports communication at the operator level, including 115.35 MB for Softmax and 53.22 MB for GELU (per Transformer block). For NEXUS, the authors report a non-interactive HE protocol requiring 164 MB of total bandwidth for BERT-base inference. For PowerFormer and the encryption-friendly architecture work, the protocols are non-interactive HE (ciphertext upload and download only), but the papers do not provide a single end-to-end inference communication size in bytes, so we do not list a numeric value for those methods. **HELIX Benchmarks.** In Table 11, we report **HELIX** in three scenarios: 1) using base embedding models (Gemini, OpenAI, Cohere) we train a linear map between each using the IMDB dataset, and report results on a downstream linear classifier for each GLUE dataset under test, 2) We use an in-distribution mapper - training the mapper using the GLUE dataset under test, and evaluating the linear classifier on the test set, and 3) We fine-tune Llama-2-8b on each dataset under test, and train a mapper between a Llama-2-8b (non-finetuned) and Llama-2-8b (fine-tuned) using the IMDB dataset.

Method	Goal	Threat	Client Inputs (Party B)	Provider Model (Party A)	Task	Interact.	Crypto
CryptoNets (ICML'16)	HE inference	HBC	✓	△	Inf.	Offline	HE
Gazelle (USENIX'18)	Fast inference	HBC	✓	✓	Inf.	Online	HE+2PC
Phoenix (CCS'22)	Reliable inference	HBC	✓	✓	Inf.	Online	HE+2PC
MPCFormer (ICLR'23)	MPC Transformer inference	HBC	✓	✓	Inf.	Online	MPC
Iron (NeurIPS'22)	Private Transformer inference	HBC	✓	✓	Inf.	Online	HE+MPC
BOLT (SP'24)	Fast private inference	HBC	✓	✓	Inf.	Online	MPC
Nimbus (NeurIPS'24)	Efficient Transformer inference	HBC	✓	✓	Inf.	Online	MPC
SecFormer (ACL'24)	Secure Transformer inference	HBC	✓	✓	Inf.	Online	MPC
HETAL (ICML'24)	HE transfer learning	HBC	✓	✗	Inf.+FT	Online	HE
Powerformer (ePrint'24)	Faster HE inference	HBC	✓	✓	Inf.	Offline	HE
NEXUS (NDSS'25)	Non-interactive secure inference	HBC	✓	✓	Inf.	Offline	HE
Enc.-Friendly LLM Architecture (ICLR'25)	HE-friendly architecture	HBC	✓	✗	Inf.+FT	Offline	HE
HELIX (ours)	Cross-model transfer	HBC	✓	✗	Inf.	Offline	HE

Table 10. Comparison of privacy-preserving ML systems. ✓ denotes cryptographic protection; △ partial protection; ✗ not protected or out of scope. **Provider Model (Party A)** indicates cryptographic protection of the provider's model parameters against disclosure to the client (black-box access permitted). **Interact.**: Offline methods rely on non-interactive HE evaluation; Online methods require MPC/2PC-style interaction. **HELIX** protects client inputs via HE but uses public/black-box access to the provider's model head, learning only a linear mapping.

1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723

Method	Model(s)	Security Scope	SST-2	STS-B (Pearson)	MRPC (F1)	RTE	Inf. Time (s)	Inf. Comm. Cost
Baseline	BERT	None	92.3	89.1	90.3	69.7	< 1	-
Iron	BERT	Full (HE+MPC)	92.8	89.4	89.9	70.8	> 60	280.99 GB
BOLT	BERT	Full (HE+MPC)	92.8	88.4	90.0	69.3	> 60	25.74 GB
SecFormer	BERT	Full (HE+MPC)	-	87.4	89.2	69.0	19	-
MPCFormer	BERT	Full (MPC)	-	80.3	88.7	64.9	18	68.6 GB
Enc.-Friendly Arch.	BERT	Full (HE)	81.9	80.0	81.5	59.3	26.5	input + output
Nimbus	BERT	Full (HE+MPC)	92.6	87.9	89.8	66.8	> 20	> 2GB
PowerFormer	BERT	Full (HE)	92.0	-	87.8	69.8	> 20	input + output
Nexus	BERT	Full (HE)	92.1	-	-	69.9	37.3	164 MB
Nexus	Llama-3-8B	Full (HE)	94.5	-	-	81.2	-	-
HELIX (IMDB Mapper, no FT)	Model Pairs*	Linear (HE)	92.3	61.0	77.8	59.6	< 1	< 1 MB
HELIX (In-dist Mapper, no FT)	Model Pairs*	Linear (HE)	92.8	75.6	80.2	61.0	< 1	< 1 MB
HELIX (Fine-Tuned)	Llama-2-8b	Linear (HE)	93.0	80.6	82.0	55.6	< 1	< 1 MB

1736
1737 *Table 11.* Accuracy and inference-time comparison on GLUE tasks. Inference time is reported per example. For HELIX, we average five
1738 runs over different random model pairs. Inference times are reported per batch. Public data mapper indicates that the model provider's
1739 data is protected as well. HELIX results are without fine-tuning. HELIX was assessed using Gemini, OpenAI, and Cohere (indicated as
1740 **Model Pairs***).

1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759

1760 J. Privacy Analysis: Membership Inference on W^*

1761 When Party A augments public mapper training with private samples (Section 5.1, Setting 2), those samples are shared in
 1762 plaintext with Party B, violating data confidentiality. However, a secondary privacy question remains: does the resulting W^*
 1763 leak **which specific samples** were included in the shared set?

1764 This matters if Party A carefully curates the shared samples—for example, sharing only non-sensitive examples while
 1765 withholding particularly private ones. If W^* encodes detectable membership signals, an adversary could infer whether a
 1766 specific sensitive sample was included, compromising Party A’s curation strategy.

1767 J.1. Attack Methodology

1768 We implement a shadow-mapper membership inference attack. Given a target sample from SST-2, we train 200 shadow
 1769 mappers (100 IN, 100 OUT) under two conditions:

- 1771 • **IN:** Mapper trained on 100K Wikipedia + 128 SST-2 samples **including** target
- 1772 • **OUT:** Mapper trained on 100K Wikipedia + 128 SST-2 samples **excluding** target

1773 For each W^* , we extract geometric features: Frobenius and spectral norms, row/column statistics, top 64 singular values,
 1774 effective rank, and bias norms. We train a logistic regression classifier to predict IN vs OUT via 5-fold stratified cross-
 1775 validation.

1776 **Configuration:** SST-2, Party A = Gemini, Party B = OpenAI, target = train index 0.

1777 J.2. Results

1778 The membership classifier achieves **0.530 ± 0.113 accuracy** (chance = 0.500), performing at chance level with substantial
 1779 variance across folds.

1780 J.3. Theoretical Hardness Analysis

1781 We provide a theoretical bound on the difficulty of membership inference from W^* .

1782 **Intuition:** When W^* is learned from N samples via ridge regression, adding or removing a single sample has only $O(1/N)$
 1783 influence on the result. With $N \approx 67,000$ samples in our experiments, any individual sample contributes less than 0.002%
 1784 to the final mapping. This tiny influence makes it fundamentally difficult to detect whether a specific sample was included.

1785 **Theoretical Bound:** For embeddings with dimensions $d_A \times d_B$, the maximum advantage any membership inference
 1786 classifier can achieve over random guessing is bounded by:

$$1791 \text{Advantage} \leq O\left(\frac{\sqrt{d_A \cdot d_B}}{N}\right) \quad (3)$$

1792 This bound follows from standard sensitivity analysis of ridge regression: the Frobenius norm difference $\|W^* - W_{-i}^*\|_F$
 1793 between a mapper trained with and without sample i is $O(1/\sqrt{N})$. Since our geometric features have $O(d_A \cdot d_B)$ dimensions,
 1794 the signal-to-noise ratio scales as $\sqrt{d_A \cdot d_B}/N$.

1795 **Empirical Validation:** For our configuration ($d_A = 1024$, $d_B = 1152$, $N = 67,000$), the theoretical bound predicts
 1796 maximum accuracy ≈ 0.516 . Our experimental result of 0.530 ± 0.113 has confidence intervals overlapping this bound,
 1797 confirming the theoretical prediction. The high variance ($\sigma = 0.113$) across folds indicates the classifier cannot reliably
 1798 distinguish IN vs OUT, consistent with the signal being at the noise floor.

1799 J.4. Interpretation

1800 The membership classifier achieves 0.530 ± 0.113 accuracy (chance = 0.500), performing near chance level with substantial
 1801 variance across folds. This negative result is theoretically expected: with $\sim 67K$ total samples, the per-sample influence on
 1802 W^* is $\sim 1.5 \times 10^{-5}$, far below the detection threshold of classifiers operating on geometric features.

1803 **Privacy implication:** Even when Party A shares limited private data (Setting 2), W^* does not leak fine-grained membership
 1804 information. The theoretical bound guarantees that membership inference advantage is $O(\sqrt{d}/N) \approx 0.016$, yielding
 1805 negligible privacy risk for large N . An adversary analyzing W^* cannot determine which specific samples Party A included
 1806 in the shared set, preserving Party A’s curation strategy.

1807 **Limitations.** This analysis assumes geometric features and tests a single target sample. Stronger adversaries with access
 1808 to many W^* samples or side information about Party A’s data distribution might achieve higher accuracy. However, the
 1809 fundamental $O(1/N)$ influence bound still applies, limiting the maximum achievable advantage.

K. Model Architecture Information

Model	Multimodal?	Q,K RMSNorm?	Embedding Dim
allenai_Olmo_3_7B_Instruct	No	No	4096
google_gemma_3_12b_it	Yes	Yes	3072
google_gemma_3_270m_it	No	Yes	1536
google_gemma_3_27b_it	Yes	Yes	4608
meta_llama_Llama_3.2_1B_Instruct	No	No	2048
meta_llama_Meta_Llama_3.8B_Instruct	No	No	4096
mistralai_Minstral_3_14B_Instruct_2512	Yes	No	5120
mistralai_Mistral_7B_Instruct_v0.2	No	No	4096
mistralai_Mistral_Nemo_Instruct_2407	No	No	5120
Qwen_Qwen2.5_0.5B_Instruct	No	No	896
Qwen_Qwen2.5_14B_Instruct	No	No	5120
Qwen_Qwen2.5_32B_Instruct	No	No	5120
Qwen_Qwen2.5_7B_Instruct	No	No	3584
Qwen_Qwen3_30B_A3B_Instruct_2507	No	Yes	2048
Qwen_Qwen3_4B_Instruct_2507	No	Yes	2560
swiss_ai_Apertus_8B_Instruct_2509	No	Yes	4096

Table 12. All models used in the text generation experiments.